

Big Data Analysis

- Part I : Introduction to Hadoop and Spark at TACC

April. 20, 2017

Dr. Weijia Xu(Manager), Dr. Ruizhu Huang
Data Mining & Statistics Group
Texas Advanced Computing Center (TACC)
The University of Texas at Austin

Big Data Training Series at TACC

- Introduction to Hadoop and Spark On Wrangler – Today
- Introduction to Scala/Spark - April 27, 2017 1:00- 4:30
 - 1:00-2:00 Introduce Scala programing
 - 2:00-3:00 Programming with spark using Scala
 - 3:00-4:00 Build spark application, unit test, submit, and running on Wrangler
 - 4:00-4:30 Hands on
- Data Analysis Using Hadoop/Spark - May 4, 2017
 - 1:00-2:00 Dataframe, SparkSQL
 - 2:00-3:00 Data analysis with MLlib and Graphx
 - 3:00-4:00 Spark streaming, advanced topic configuration/ optimization
 - 4:00-4:30 Hands on

Today's Agenda

1:00 – 1:30 Overview of Big Data Processing

- Software in Big Data/TACC
- Big Data and HPC
- Intro. to MapReduce Model

1:30 – 2:30 What's Hadoop

- What are HDFS and YARN
- Hadoop on Wrangler

2:30 – 3:00 Breaks / Hands-on exercise

3:00 – 4:00 Programming with Hadoop and Spark

- Using Hadoop with other languages
- Hadoop vs. Spark
- Intro. of Zeppelin
- Using Spark with Scala, Python and R

4:00 – 4:30 Hands-on exercise

What's the big challenge with the big data analysis?

Data analysis process requires even more computational resources

Storage, triple the size of the raw data to store the intermediate files, output etc.

Memory, e.g. algorithm want to store the pair-wise distance matrix among data points

The analysis process would take much longer

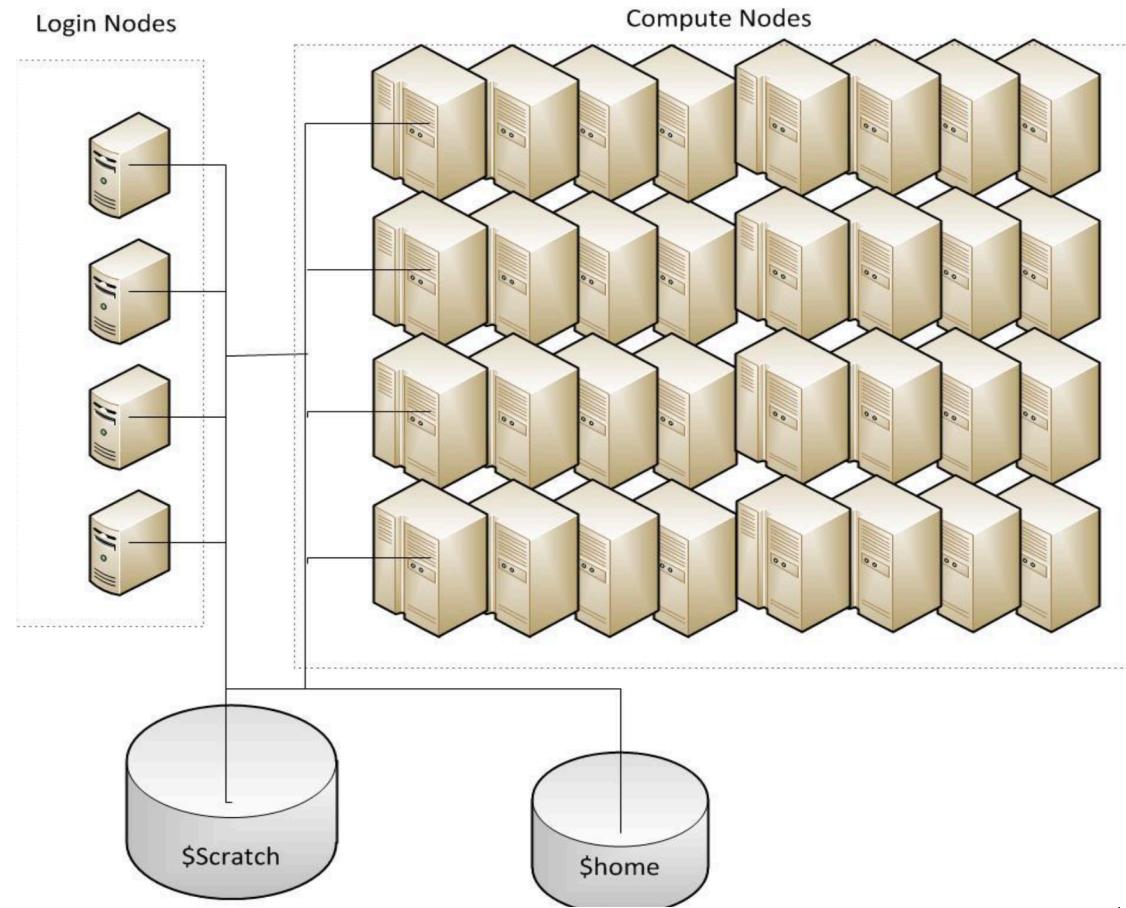
Typical hard drive read speed is about 150MB/sec,

Reading 1TB ~ 2 hours

Analysis could require processing time quadratic to the size of the data

Analysis took 1 second for 1GB data, would require 11 days to finish for 1TB data

Big Data v.s. High Performance Computing



Traditional Computation at Computing Cluster

Data are stored separately and accessed by compute nodes during the access,

Typically a high speed connection between the storage and computing resources, e.g. 40Gbps

User need develop code that can run on multiple nodes concurrently

Typically, master-slaves mode, where a master thread to control the process of multiple workers.

Do not really solve the Big Data problem.

MapReduce Model for Big Data

A programing model proposed and used by Google.

Serving as a platform for customized computation over large scale of data. :

At the core, users implement interface of two main functions:

```
map  (in_key, in_value) ->  
      (out_key, intermediate_value) list  
reduce (out_key, intermediate_value list) ->  
       out_value list
```

It's not a brand new idea, but several old ideas put together to handle data at scale

“Big” Ideas in MapReduce

Move computations to data.

Do not use/assume high bandwidth inter-connection between nodes.

Avoid and reduce the need of data transfer over the network is often the bottleneck.

Processing data in sequential order, avoid random access

Same idea as applied in the RDBMS

However, assume data can be processed any order,

All data are packed in to blocks (default at 64MB).

“Big” Ideas in MapReduce

Scale out instead of scale up:

The same code can run with 1 node or 100 nodes.

Hide system details from user

Providing abstraction in writing parallel code.

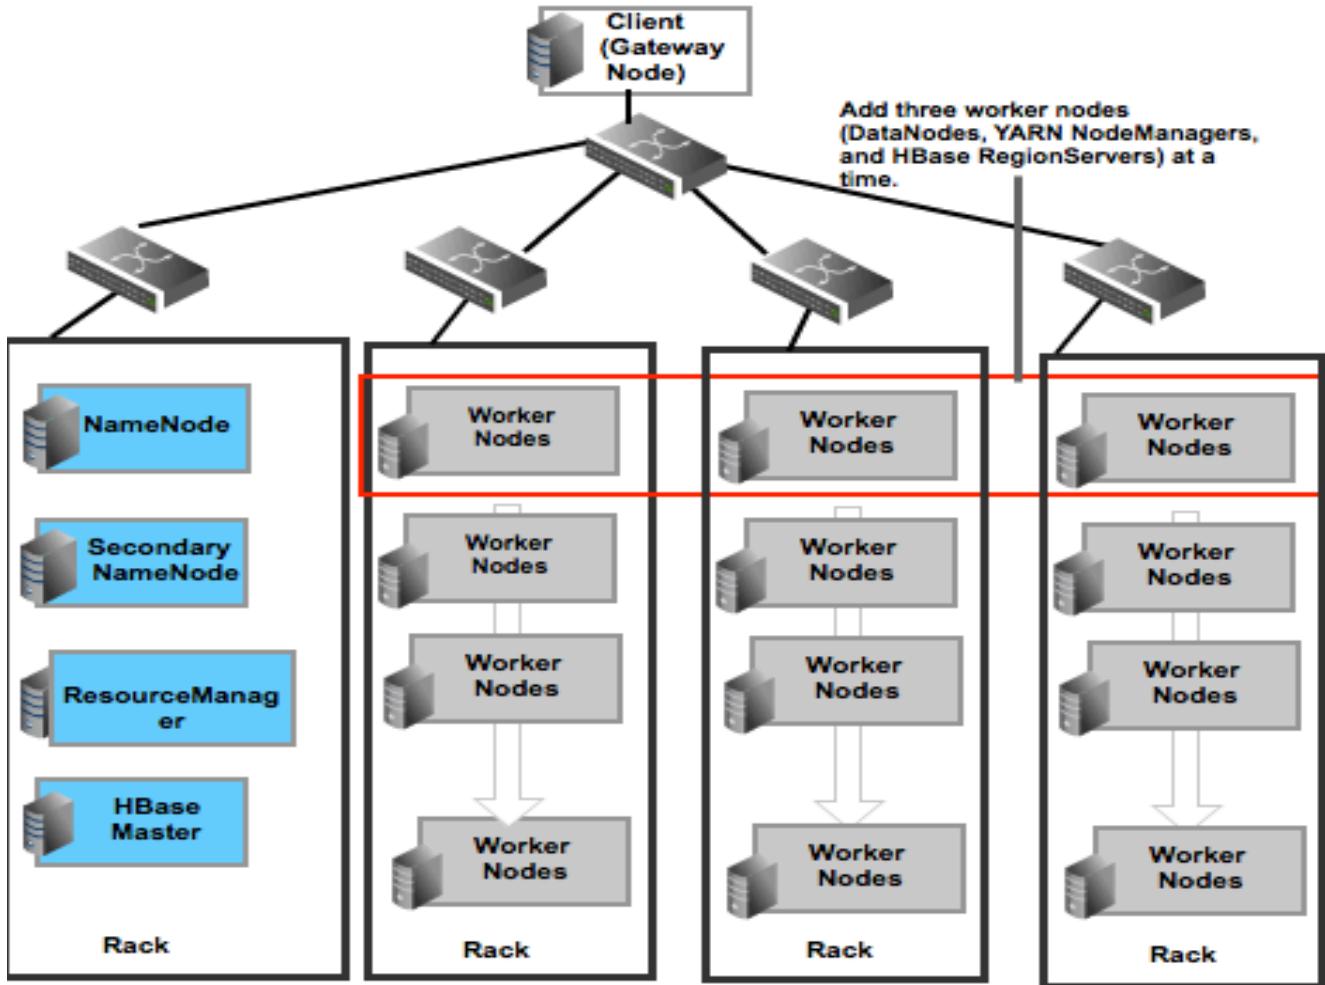
- Mapper and reducer

- Partitioner and combiner

Isolates developer from (and independent from) system hardware details

Once all required components are specified, data is automatically “sliced” and processed in parallel.

Support MapReduce with Hadoop Cluster



Support MapReduce with Hadoop Cluster

Data are stored distributed across nodes within the hadoop cluster

The workload distribution are automatically handled by the Hadoop cluster

Data can be treated as key-value pair

Either the “Key” or “Value” in a key-value pair could be any type of data.

The computation is broke down in to two major steps:

Map instances:

process the data stored within one data block sequentially,

The result is in the format of list of Key-value pairs

Reduce instances:

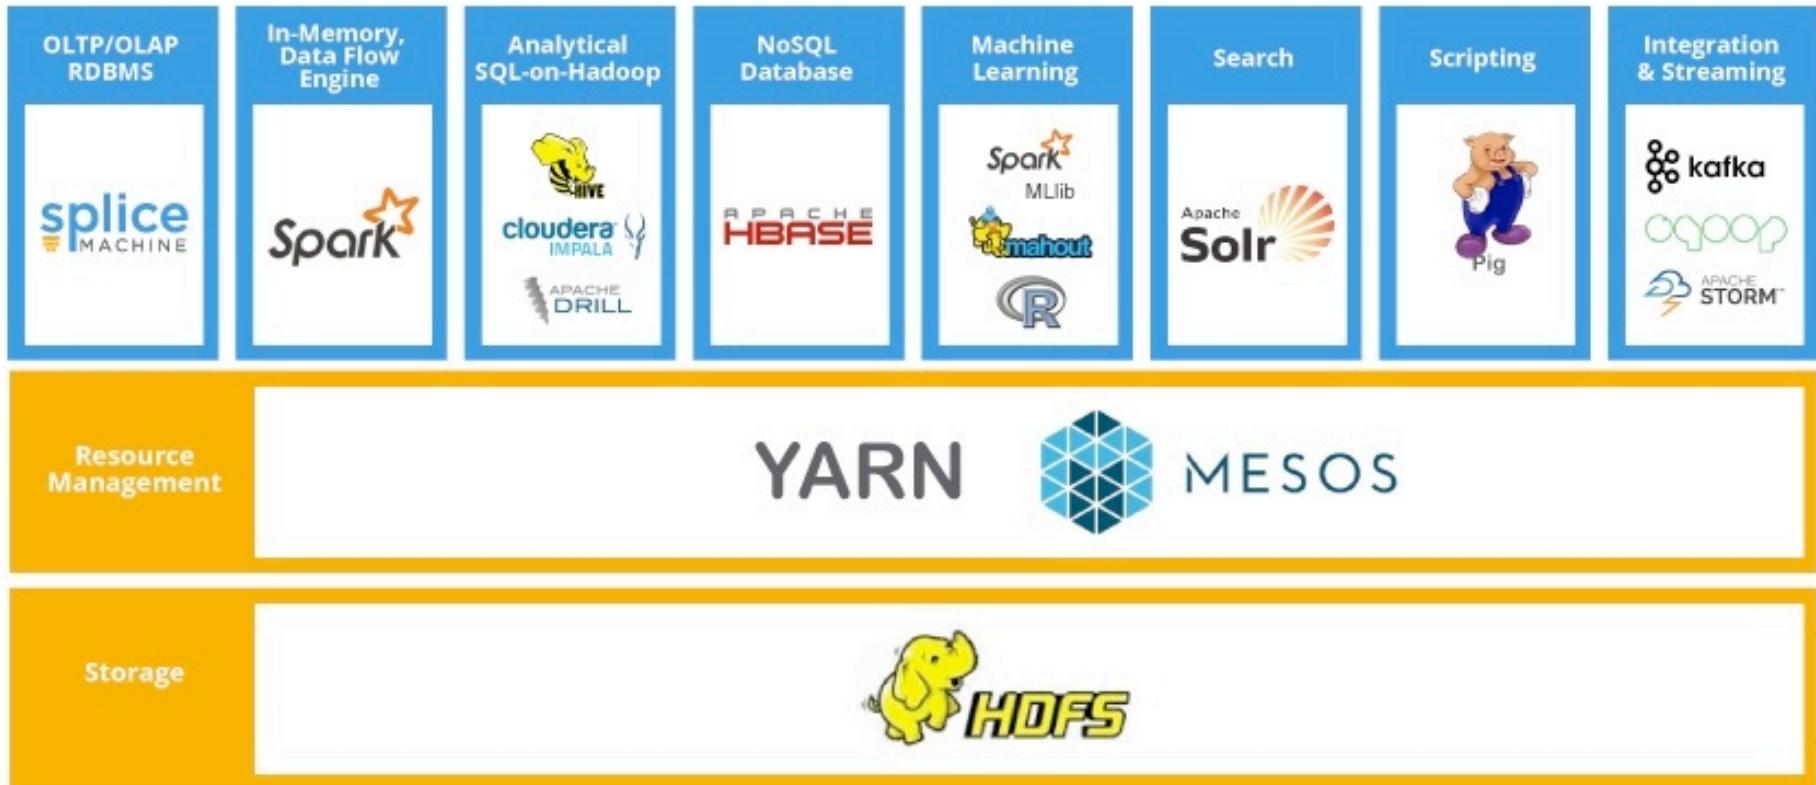
Collect (key, value) pairs emitted by Map instances

Pairs with the same key are sent to the same reducer,

Each Reducer process the key-value pairs received and write output to a file.

User only required to develop functions for Map and Reduce

Hadoop Ecosystem



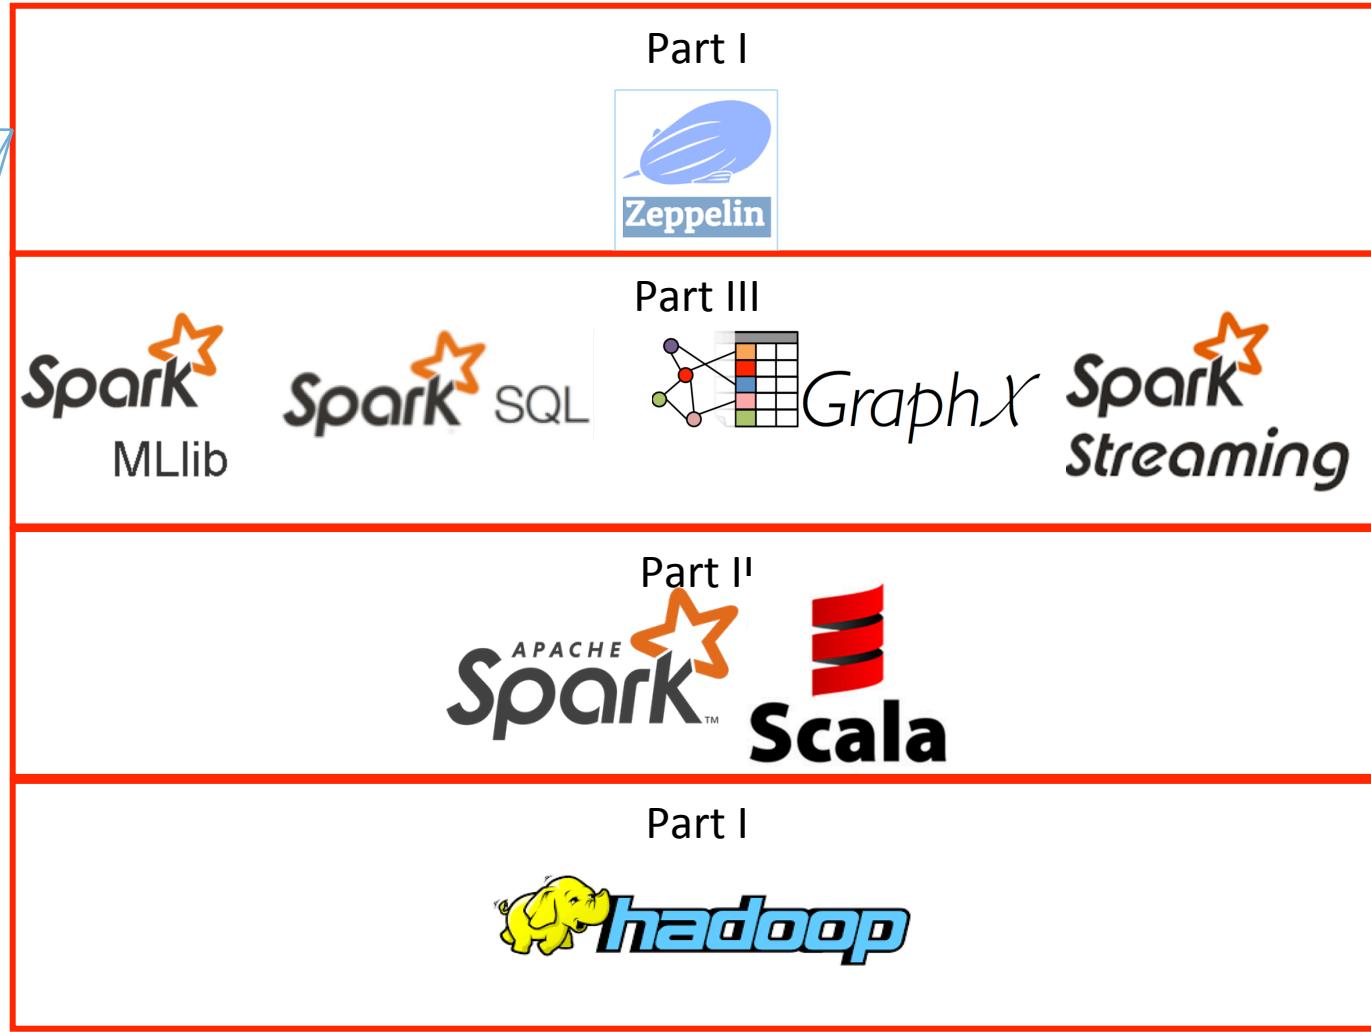
Big Data Analysis Support at TACC



Caffe



About this series

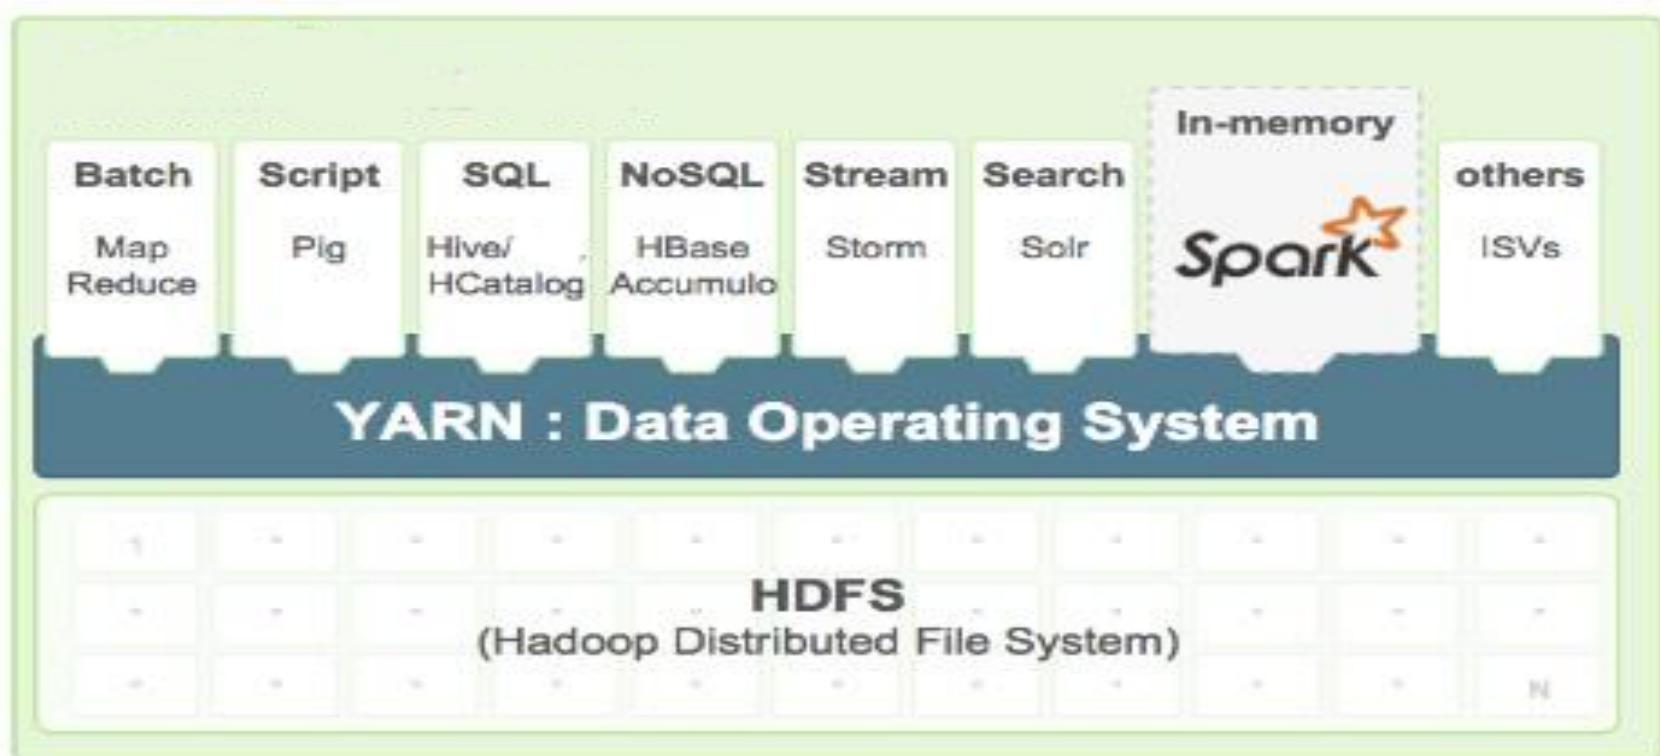


Recap

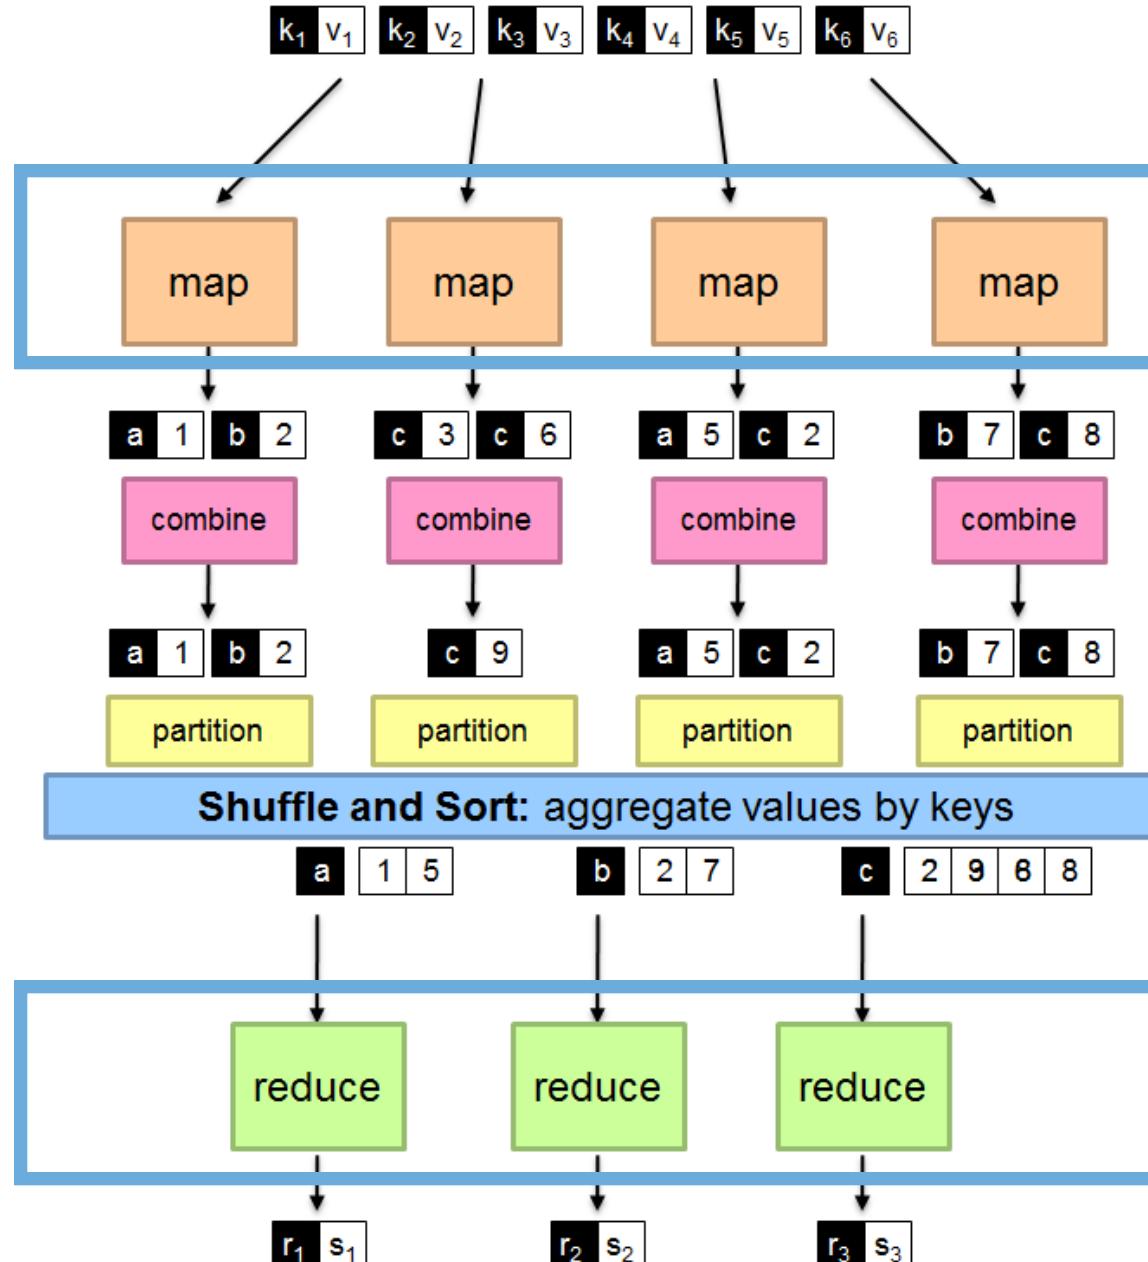
- Big Data Challenges
 - storing capacity
 - processing power
- MapReduce Model
 - Move computations to data
 - Processing data in sequential order, avoid random access
- Hadoop Infrastructure
 - NameNode
 - Secondary NameNode
 - Resource Manager
 - Worker/Compute/Data Nodes
- Software in Big Data/TACC

Hadoop and Yarn

- Hadoop is an open source implementation of MapReduce programming model in JAVA with interface to other programming language such as C/C++, python.
- Hadoop includes HDFS, YARN, MapReduce



Map Reduce Model



Word Count Example

Read text files and count how often words occur.

The input is text files

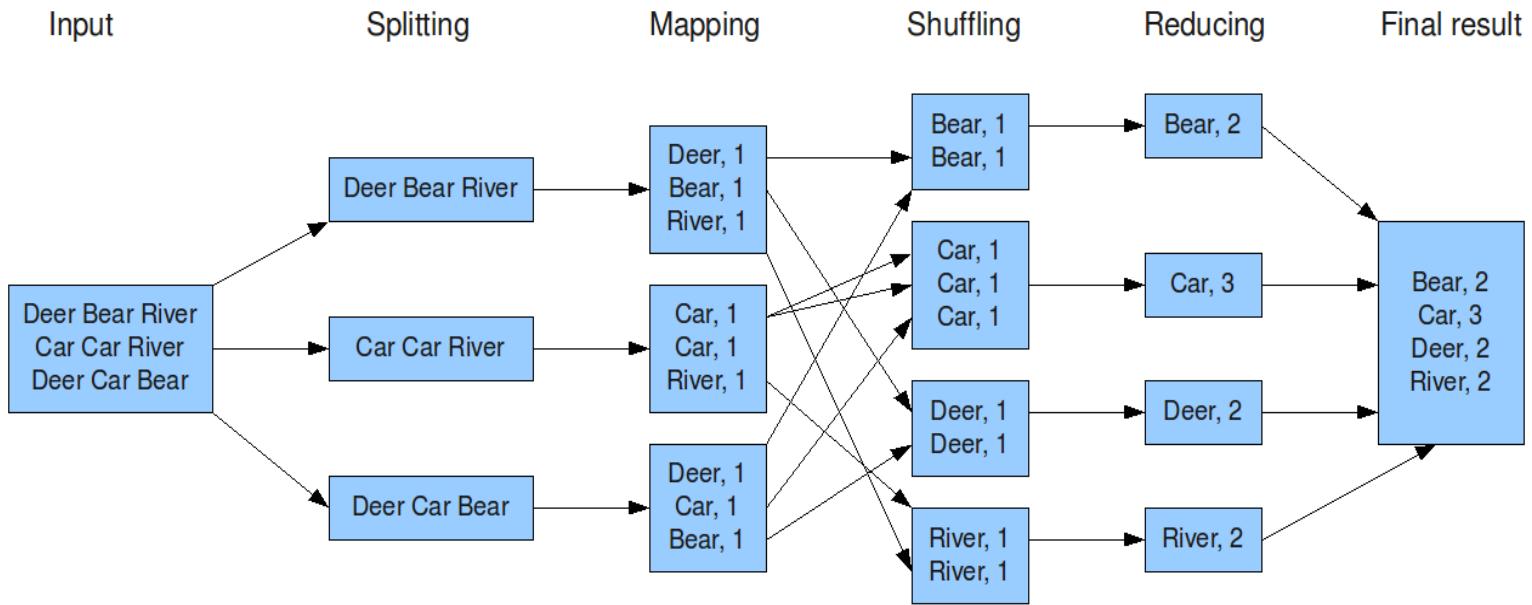
The output is a text file

each line: word, tab, count

Map: Produce pairs of (word, count)

Reduce: For each word, sum up the counts.

The overall MapReduce word count process



WordCount-Java Overview

```
3 import ...
12 public class WordCount {
14     public static class Map extends Mapper<Object, Text, Text, IntWritable> {
18         public void map ...
26     }
27
28     public static class Reduce extends Reducer<Text,IntWritable,Text,IntWritable> {
30         public void reduce ...
37     }
38
39     public static void main(String[] args) throws Exception {
40         Job job = Job.getInstance(new Configuration(), "word count");
41         ...
53         FileInputFormat.setInputPaths(conf, new Path(args[0]));
54         FileOutputFormat.setOutputPath(conf, new Path(args[1]));
55
56         System.exit(job.waitForCompletion(true) ? 0 : 1);
57     }
59 }
```

wordCount Mapper - Java

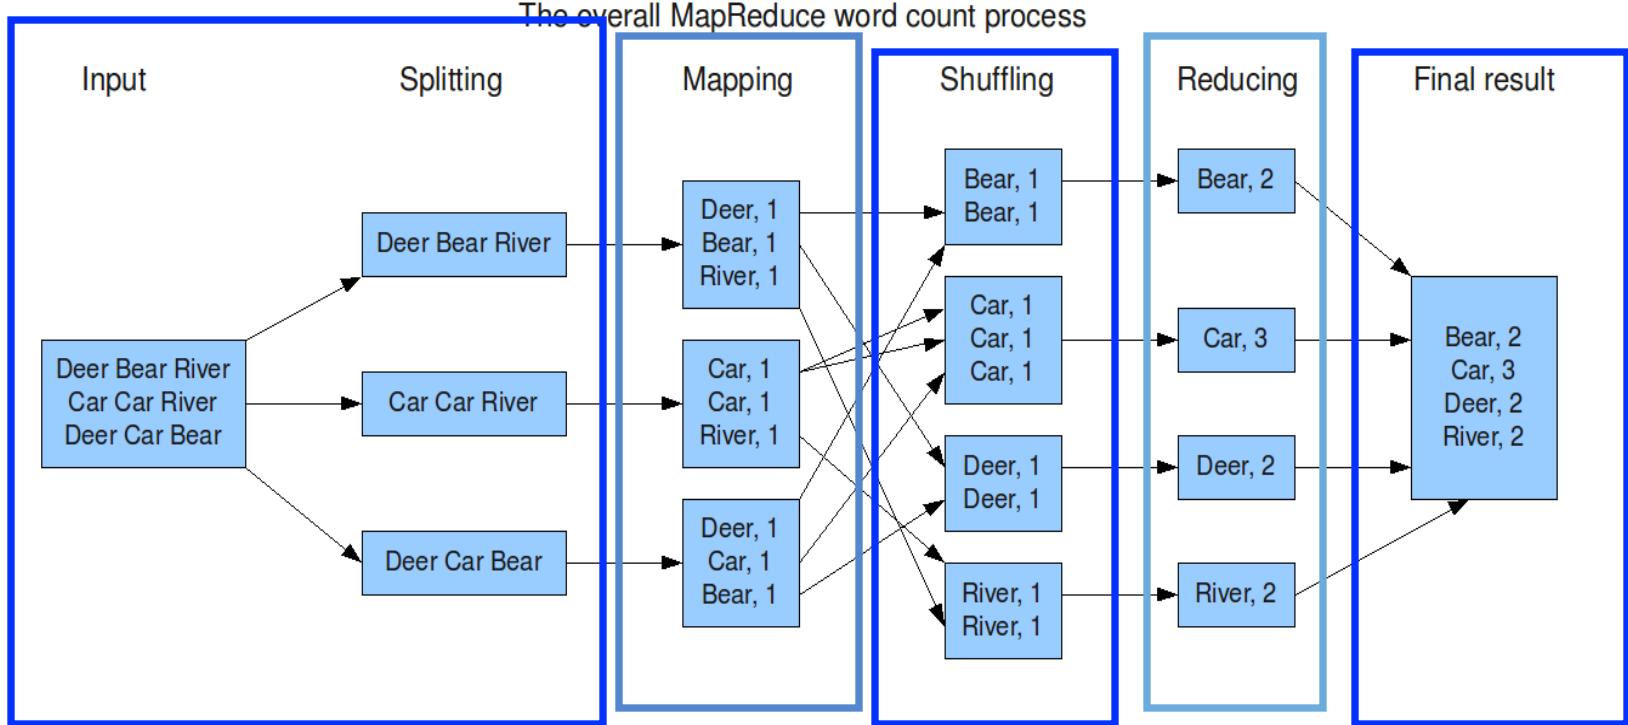
```
14 public static class Map extends Mapper<Object, Text, Text, IntWritable> {  
15     private final static IntWritable one = new IntWritable(1);  
16     private Text word = new Text();  
17  
18     public void map( Object key, Text value, Context context)  
         throws IOException {  
19             String line = value.toString();  
20             StringTokenizer tokenizer = new  
StringTokenizer(line);  
21             while (tokenizer.hasMoreTokens()) {  
22                 word.set(tokenizer.nextToken());  
23                 context.write(word, one);  
24             }  
25     }  
26 }
```

wordCount Reducer - Java

```
28 public static class Reduce extends Reducer<Text, IntWritable,  
Text, IntWritable> {  
29  
30     public void reduce(Text key, Iterable<IntWritable> values,  
                         Context context)  
throws IOException {  
31         int sum = 0;  
32         while (values.hasNext()) {  
33             sum += values.next().get();  
34         }  
35         result.set(sum);  
36         context.write(key, result);  
37     }  
38 }
```

WordCount main - Java

```
public static void main(String[] args) throws Exception {  
    Job job = Job.getInstance(new Configuration(), "word count");  
  
    job.setJarByClass(WordCount.class);  
    job.setMapperClass(Map.class);  
    job.setCombinerClass(Reduce.class);  
    job.setReducerClass(Reduce.class);  
    job.setOutputKeyClass(Text.class);  
    job.setOutputValueClass(IntWritable.class);  
  
    FileInputFormat.addInputPath(job, new Path(args[0]));  
    FileOutputFormat.setOutputPath(job, new Path(args[1]));  
    System.exit(job.waitForCompletion(true) ? 0 : 1);  
}
```



Mapper - Java

Maps **input key-value** pairs to a set of **intermediate key-value** pair

Class for Individual tasks to run.

One mapper task per *InputSplit*.

Map function are automatically called per key Value pair.

```
public static class Map extends Mapper<Object, Text, Text, IntWritable> {  
    private final static IntWritable one = new IntWritable(1);  
    private Text word = new Text();  
  
    public void map( Object key, Text value, Context context)  
        throws IOException {  
        String line = value.toString();  
        StringTokenizer tokenizer = new StringTokenizer(line);  
        while (tokenizer.hasMoreTokens()) {  
            word.set(tokenizer.nextToken());  
            context.write(word, one);  
        }  
    }  
}
```

Reducer - Java

Reduces the **set of values** of the **same key** to a **smaller set**.

Each reducer will process subset generated by partitioner

Each reducer will generate an output file

```
public static class Reduce extends Reducer<Text, IntWritable, Text,  
    IntWritable> {  
    public void reduce(Text key, Iterable<IntWritable> values,  
        Context context) throws IOException {  
        int sum = 0;  
        while (values.hasNext()) {  
            sum += values.next().get();  
        }  
        result.set(sum);  
        context.write(key, result);  
    }  
}
```

Mapper - Python

```
# input comes from STDIN (standard input)
for line in sys.stdin:
    # remove leading and trailing whitespace
    line = line.strip()
    # split the line into words
    words = line.split()
    # increase counters
    for word in words:
        # write the results to STDOUT (standard output);
        # what we output here will be the input for the
        # Reduce step, i.e. the input for reducer.py
        #
        # tab-delimited; the trivial word count is 1
        print '%s\t%s' % (word, 1)
```

Reducer - Python

```
# input comes from STDIN
for line in sys.stdin:
    # remove leading and trailing whitespace
    line = line.strip()

    # parse the input we got from mapper.py
    word, count = line.split('\t', 1)

    # convert count (currently a string) to int
    try:
        count = int(count)
    except ValueError:
        # count was not a number, so silently
        # ignore/discard this line
        continue
```

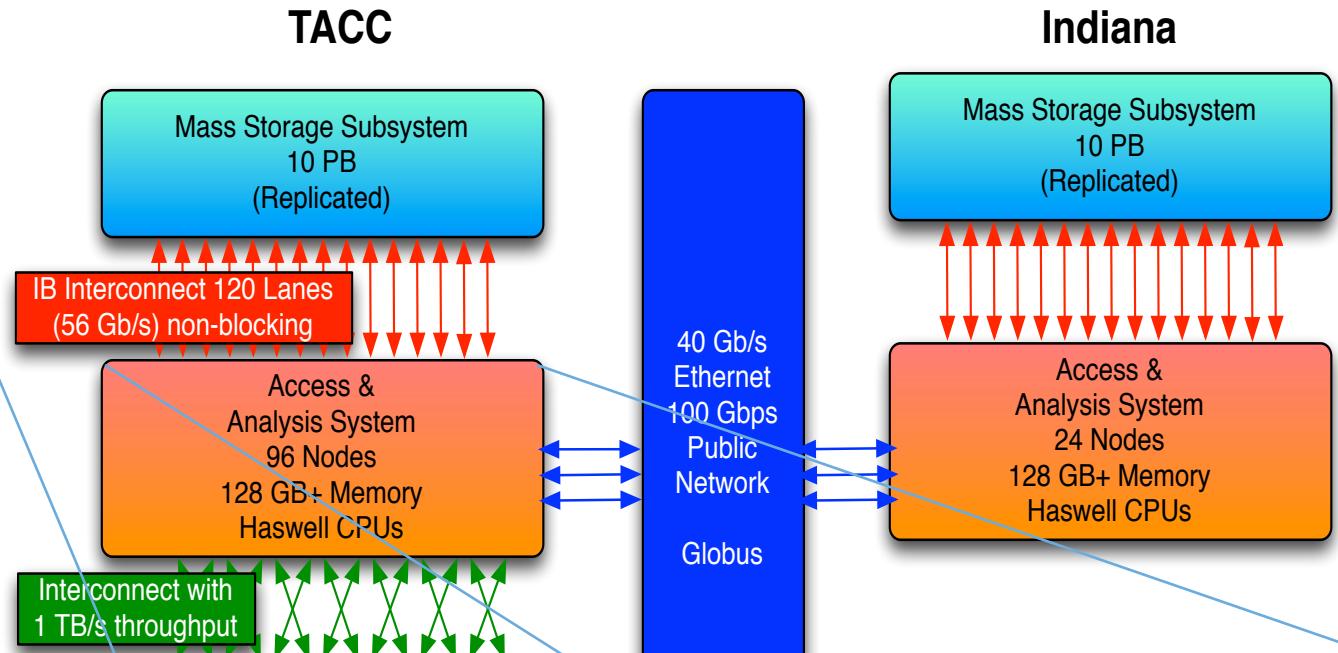
```
# this IF-switch only works because Hadoop
# sorts map output
# by key (here: word) before it is passed to
# the reducer
if current_word == word:
    current_count += count
else:
    if current_word:
        # write result to STDOUT
        print '%s\t%s' % (current_word,
                           current_count)
    current_count = count
    current_word = word

# do not forget to output the last word if
# needed!
if current_word == word:
    print '%s\t%s' % (current_word,
                       current_count)
```

Accessing Hadoop Cluster on Wrangler

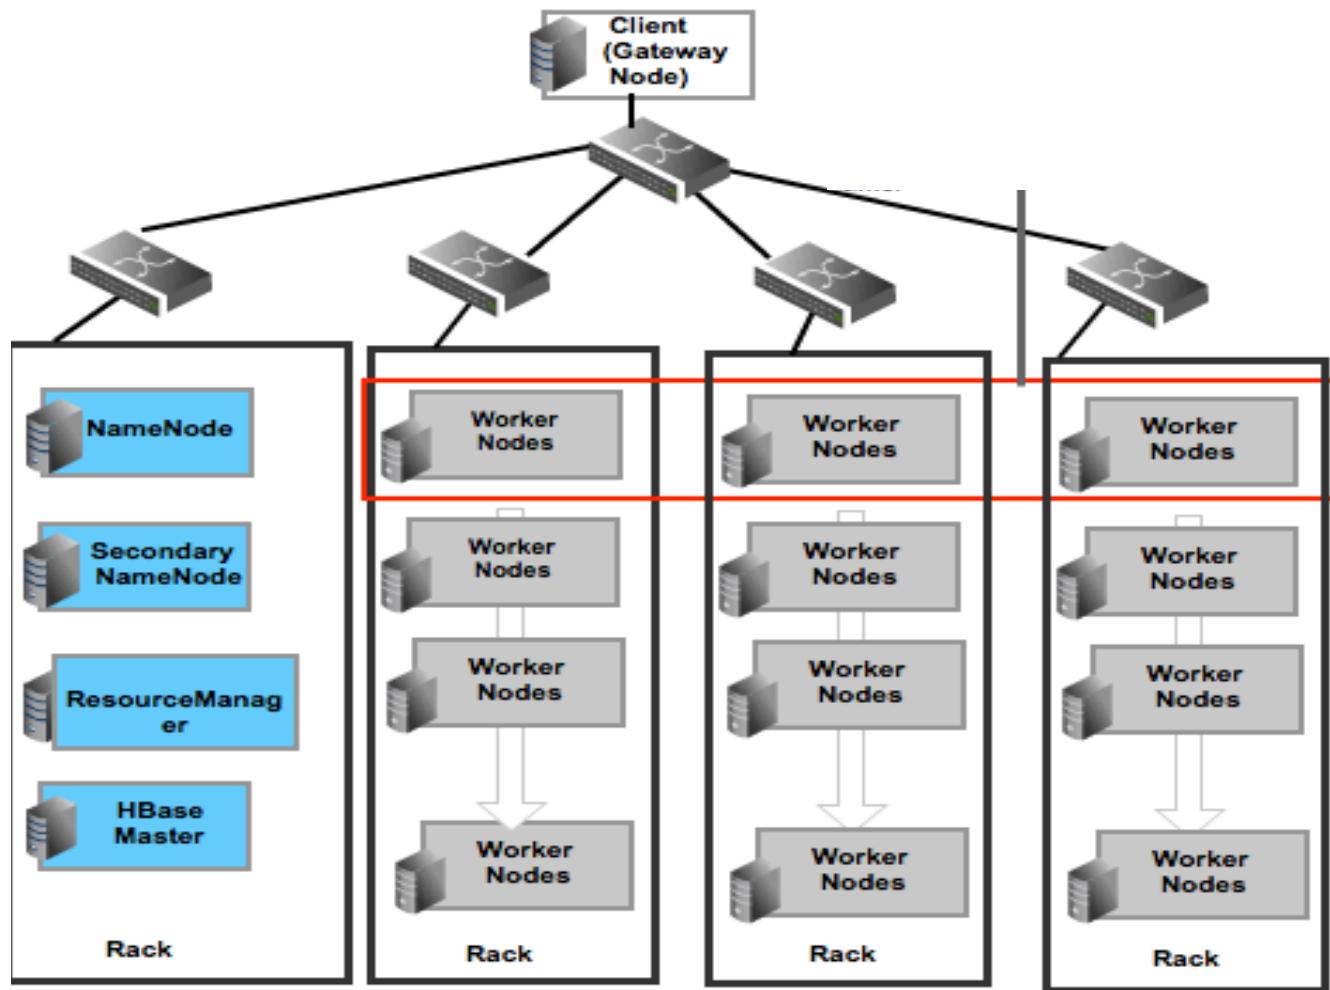
Wrangler

- A direct attached PCI interface allows access to the NAND flash.
- Not limited by networking connection
- Flash storage not tied to individual nodes



- The Hadoop cluster can be dynamically created over 2 to 48 nodes for each project to use in allocated time
- Each node has access to 4 TB flash storage across four channels
- Accessible via the Hadoop cluster via idev, batch job submission and VNC sessions.

A Typical Hadoop Infrastructure



Get Started with Hadoop on Wrangler

Step 1: create a Hadoop reservation through Wrangler data portal

What do you need?

Any browser

Done !!!

Reservation Name: hadoop+TRAINING-HPC+2186

Step 2: Access your Hadoop reservation to run Hadoop jobs

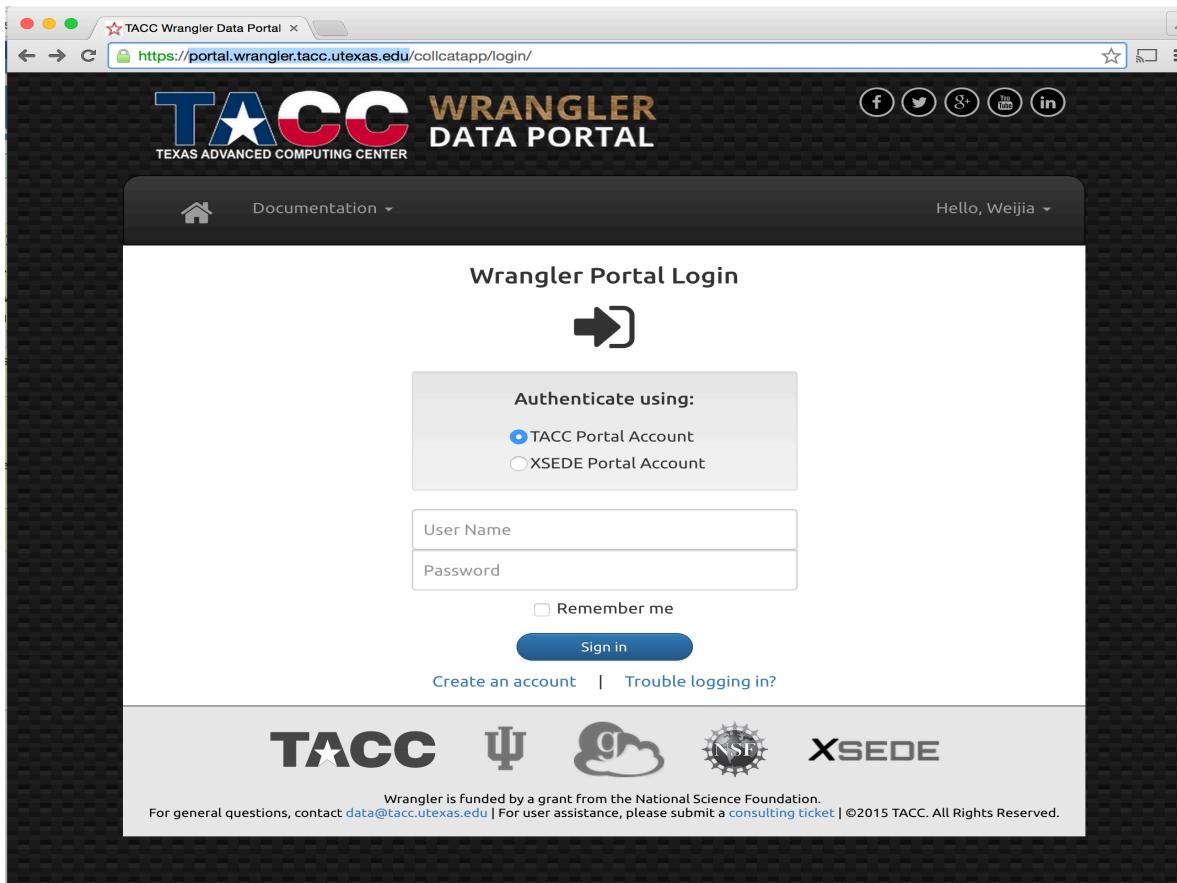
What do you need?

Secure Shell Client

Any VNC client

Create Hadoop Reservation

Wrangler data portal: [portal.wrangler.tacc.utexas.edu](https://portal.wrangler.tacc.utexas.edu/collcatapp/login/)



On project page choose: Manage -> Create Hadoop Reservation

Hadoop Reservations

Overview

Hadoop Reservations result in the automatic creation of a Hadoop Cluster on a set of D5 nodes. Hadoop Reservations are charged against the Compute NUs portion of your award. The number of nodes and subsequent cost of the reservation are determined by the amount of data to be managed within the created Hadoop File System (HDFS). By default, HDFS is setup to use a replication factor of 2. You have the option of altering the replication factor after the reservation begins and the HDFS file system has been created for you.

If you need to create a Hadoop Reservation for more than 10 total nodes, please create a [consulting ticket](#) and TACC Staff will work with you on these special requests.

Calculating Allocation Costs

Example: Your project requires 15 TB total usable space for data input, intermediary and output. You want this data campaign to last 10 days.

Step	Calculation	Cost
1	15TB x 2 replication factor	30TB
2	30TB / 4TB per data node	8 data nodes required
3	8 data nodes + 1 name node	9 total nodes
4	9 nodes x 10 days x 24 hrs	2,160 NUs required
TOTAL		2,160 NUs

Current NU Balances

Awarded	Charged	Pending	Balance
5000	0	0	5000

Request New Hadoop Reservation

Number of Nodes*
(2-10)

Duration*
Days (1-30)

Start as soon as possible? Date Selection
OR Schedule a start date

TACC    **XSEDE**

Wrangler is funded by a grant from the National Science Foundation.
For general questions, contact data@tacc.utexas.edu | For user assistance, please submit a [consulting ticket](#) | ©2015 TACC. All Rights Reserved.

the number of nodes (1 ~10) to be used for the Hadoop cluster.

Duration (1-30 Days)

Schedule Start time

Check Hadoop Reservation

- Once log on to Wrangler login node, user can check the reservation status with `scontrol` command:

>scontrol show reservation

```
login1.wrangler(9)$ scontrol show reservation
ReservationName=Service_nodes_equivalent StartTime=2017-03-28T12:59:00 EndTime=2017-04-25T13:00:00 Duration=28-00:01:00
  Nodes=c251-[140-143] NodeCnt=4 CoreCnt=192 Features=(null) PartitionName=(null) Flags=MAINT,OVERLAP,IGNORE_JOBS,SPEC_NODES
  Users=root,nthorne,plubbs,ctjordan,ngaffney,xwj Accounts=(null) Licenses=(null) State=ACTIVE

ReservationName=hadoop+PEMFS+2178 StartTime=2017-04-15T10:20:02 EndTime=2017-04-20T10:20:02 Duration=5-00:00:00
  Nodes=c252-[109-112] NodeCnt=4 CoreCnt=192 Features=(null) PartitionName=hadoop Flags=STATIC
  Users=gruan,xwj,agupta,rhuang,avliu,ychu,zhang,parnell Accounts=(null) Licenses=(null) State=ACTIVE
```

- The reservation will include all users from the projects
- The first node in the reservation will be used as ***namenode***
- The hadoop cluster will start with a set of default settings
 - User may override most settings such as duplication factor, block size at run time per application.
 - Hadoop cluster with specific settings upon request

Access Hadoop Reservation

Once the reservation status is “active”, a user can access through slurm job:

VNC job: starts a vnc server session on one of the node in Hadoop cluster,

- Check cluster information and hadoop job status
- Application with Graphical/Web user interface

idev job: Assign one node in Hadoop cluster to user

- Manage data in and out hadoop cluster,
- Submit Hadoop jobs via command line
- Code testing

Batch job: submitting jobs to YARN resource manager in Hadoop cluster.

- Submit large analysis job
- Submit batch of processing jobs to run sequentially

Access Hadoop Cluster with VNC

Please visit: vis.tacc.utexas.edu

The screenshot shows a web browser window for the TACC Visualization Portal at <https://vis.tacc.utexas.edu/>. The page title is "TACC Visualization Portal". At the top right, it says "Not logged in.". Below the title, there are three navigation links: "Home", "Jobs", and "Help". A large blue callout points from the text "Choose ‘TACC User Portal User’" to the radio button labeled "TACC User Portal User", which is selected. Another red oval highlights the "Username" and "Password" input fields. A blue line connects the "Login" button to the text "Enter credential". At the bottom of the page, there are two smaller windows showing "Job Submission" and "VNC Visualization Session".

Welcome to the TACC Visualization Portal
Simple access to TACC's Vis Resources

Features:

- Remote, interactive, web-based visualization
- iPython / Jupyter Notebook integration
- R Studio integration
- Run on [Maverick](#) and [Stampede](#)
- Visualization job submission and monitoring
- Current resource usage and allocation view

Job Submission VNC Visualization Session

Choose “TACC User Portal User”

Enter credential

https://vis.tacc.utexas.edu/#

TACC Visualization Portal

TACC\train221 logout
No job running.

Home Jobs Help

Start a Job

Resource  Maverick Stampede Wrangler

Project  TRAINING-HPC

Session type VNC
 iPython/Jupyter Notebook
 R Studio

Reservation ID  Hadoop+SC15

Queue  hadoop

Desktop resolution 1920x1080

Number of nodes 1

Wayness (processes per node) 20

Note: increasing the number of nodes will only increase performance for parallel applications (e.g. ParaView or VisIt). The wayness parameter is only relevant to parallel applications, and determines how many processes are spawned per node when the parallel application is executed.

Start Job 

Set VNC Password 

1. Choose Wrangler Tab
2. Set VNC password, (Only need once)
3. Fill in reservation name:
hadoop+TRAINING-HPC+2186
And choose "hadoop" queue

c252-102.wrangler.tacc.utexas.edu:1 (xwj)

Applications Places System

*** Exit this window to kill your VNC server ***

Wed Apr 22, 2:15 PM Weijia Xu

Namenode Information - Mozilla Firefox

MapReduce Job job_14297... x Namenode Information x All Applications x +

c252-101.wrangler.tacc.utexas.edu:50070/dfshealth.html#tab-overview

8 Google

Hadoop Overview Datanodes Snapshot Startup Progress Utilities

Overview 'c252-101.wrangler.tacc.utexas.edu:8020' (active)

Started:	Wed Apr 22 12:19:26 CDT 2015
Version:	2.5.0-cdh5.3.0, r19097cda2536da1d41ff6713556c8f7284174d
Compiled:	2014-12-17T03:05Z by jenkins from Unknown
Cluster ID:	CID-dda5fe29-0ee3-4a1b-946e-298b6c57bb63
Block Pool ID:	BP-1405723930-129.114.58.144-1429723164687

Summary

Security is off.

Safemode is off.

2504 files and directories, 9032 blocks = 11536 total filesystem object(s).

Heap Memory used 175.5 MB of 893 MB Heap Memory. Max Heap Memory is 893 MB.

Non Heap Memory used 34.8 MB of 35.31 MB Committed Non Heap Memory. Max Non Heap Memory is 214 MB.

Configured Capacity:	26.91 TB
DFS Used:	3.76 TB
Non DFS Used:	2.52 TB
DFS Remaining:	20.63 TB
DFS Used%:	13.97%
DFS Remaining%:	76.03%

[VNC config] *** Exit this window to... Namenode information...



Applications Places System



*** Exit this window to kill your VNC server ***

Namenode information - Mozilla Firefox

http://c252-1...5734_0012/xwj x

Namenode information

x All Applications

x +

c252-101.wrangler.tacc.utexas.edu:50070/dfshealth.html#tab-overview

g v Google



Summary

Security is off.

Safemode is off.

2504 files and directories, 9032 blocks = 11536 total filesystem object(s).

Heap Memory used 175.5 MB of 893 MB Heap Memory. Max Heap Memory is 893 MB.

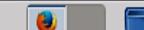
Non Heap Memory used 34.8 MB of 35.31 MB Committed Non Heap Memory. Max Non Heap Memory is 214 MB.

Configured Capacity:	26.91 TB
DFS Used:	3.76 TB
Non DFS Used:	2.52 TB
DFS Remaining:	20.63 TB
DFS Used%:	13.97%
DFS Remaining%:	76.65%
Block Pool Used:	3.76 TB
Block Pool Used%:	13.97%
DataNodes usages% (Min/Median/Max/stdDev):	13.49% / 13.95% / 14.65% / 0.30%
Live Nodes	14 (Decommissioned: 0)
Dead Nodes	0 (Decommissioned: 0)
Decommissioning Nodes	0
Number of Under-Replicated Blocks	0
Number of Blocks Pending Deletion	0

[VNC config]

*** Exit this window to...

Namenode information...



c252-102.wrangler.tacc.utexas.edu:1 (xw)

Applications Places System

Browse and run installed applications *** Exit this window to kill your VNC server ***

All Applications - Mozilla Firefox

MapReduce Job job_14297... x | Namenode information x | All Applications x +

c252-101.wrangler.tacc.utexas.edu:8088/cluster/apps g Google

 All Applications

Logged in as: dr.who

All Applications

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved	Active Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes
14	0	1	13	406	1.58 TB	1.60 TB	56 GB	406	672	14	14	0	0	0	0

User Metrics for dr.who

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Containers Pending	Containers Reserved	Memory Used	Memory Pending	Memory Reserved	VCores Used	VCores Pending	VCores Reserved
0	0	1	13	0	0	0	0 B	0 B	0 B	0	0	0

Show 20 entries Search:

ID	User	Name	Application Type	Queue	StartTime	FinishTime	State	FinalStatus	Progress	Tracking UI
application_1429723175734_0014	xwj	TeraGen	MAPREDUCE	root.xwj	Wed Apr 22 14:05:23 -0500 2015	N/A	RUNNING	UNDEFINED	<div style="width: 0%;"></div>	ApplicationMaster
application_1429723175734_0013	xwj	TeraSort	MAPREDUCE	root.xwj	Wed Apr 22 12:57:56 -0500 2015	Wed Apr 22 13:06:04 -0500 2015	FINISHED	SUCCEEDED	<div style="width: 100%;"></div>	History
application_1429723175734_0012	xwj	TeraGen	MAPREDUCE	root.xwj	Wed Apr 22 12:54:20 -0500 2015	Wed Apr 22 12:56:46 -0500 2015	FINISHED	SUCCEEDED	<div style="width: 100%;"></div>	History
application_1429723175734_0011	xwj	hadoop-mapreduce-client-jobclient-2.5.0-cdh5.3.0-tests.jar	MAPREDUCE	root.xwj	Wed Apr 22 12:51:42 -0500 2015	Wed Apr 22 12:52:29 -0500 2015	FINISHED	SUCCEEDED	<div style="width: 100%;"></div>	History
application_1429723175734_0010	xwj	hadoop-mapreduce-client-jobclient-2.5.0-cdh5.3.0-tests.jar	MAPREDUCE	root.xwj	Wed Apr 22 12:50:36 -0500 2015	Wed Apr 22 12:51:22 -0500 2015	FINISHED	SUCCEEDED	<div style="width: 100%;"></div>	History
application_1429723175734_0009	xwj	hadoop-	MAPREDUCE	root.xwj	Wed Apr 22	Wed Apr 22	FINISHED	SUCCEEDED	<div style="width: 100%;"></div>	History

[VNC config] *** Exit this window to... All Applications - Mozi...

c252-102.wrangler.tacc.utexas.edu:1 (xw)

Applications Places System

Browse and run installed applications *** Exit this window to kill your VNC server

MapReduce Job job_1429723175734_0012 - Mozilla Firefox

MapReduce Job job_1429723175734_0012 x Namenode Information x All Applications x +

c252-101.wrangler.tacc.utexas.edu:19888/jobhistory/job/job_1429723175734_0012 v C Google

Logged in as: dr.who

 MapReduce Job job_1429723175734_0012

Job Overview

Application

+ Job

- Overview
- Counters
- Configuration
- Map tasks
- Reduce tasks

+ Tools

Job Name: TeraGen
User Name: xwj
Queue: root.xwj
State: SUCCEEDED
Uberized: false
Submitted: Wed Apr 22 12:54:20 CDT 2015
Started: Wed Apr 22 12:54:24 CDT 2015
Finished: Wed Apr 22 12:56:46 CDT 2015
Elapsed: 2mins, 22sec

Diagnostics:

Average Map Time 1mins, 41sec

ApplicationMaster

Attempt Number	Start Time	Node	Logs
1	Wed Apr 22 12:54:21 CDT 2015	c252-113.wrangler.tacc.utexas.edu:8042	logs

Task Type

Map	Total	Complete
256	256	256
Reduce	0	0

Attempt Type

Maps	Failed	Killed	Successful
0	2	256	256
Reduces	0	0	0

About Apache Hadoop

[VNC config] *** Exit this window to... MapReduce Job job_1...

42

Access Hadoop Reservation via idev Session

User can submit idev session to hadoop cluster reservation

➤ ***idev -r*** hadoop+TRAINING-HPC+2186

It defaults to use your default project,

The ***-A allocation_name*** option to specify allocation to use

The default duration for idev is 30 minutes

The ***-m minutes*** option can specify the time of the idev session

Please limit your usage to Hadoop related tasks, you can also submit idev without using reservation for non-hadoop tasks.

Access by secure shell client

```
ssh username@wrangler.tacc.utexas.edu
```

```
idev -r hadoop+TRAINING-HPC+2186 -m 240 -p hadoop
```

- Access by vis portal
 - Go to vis.tacc.utexas.edu using web browser
 - Login with your credential
 - Goto Wrangler tab to start VNC sessions using reservation hadoop+TRAINING-HPC+2186 **and using Hadoop queue**

Hadoop Distributed File System (HDFS)

The hdfs will be set up with three top level directories

```
c252-101.wrangler(3)$ hadoop fs -ls /
Found 3 items
drwxrwxrwx  - hdfs hadoop      0 2015-06-11 18:49 /tmp
drwxr-xr-x  - hdfs hadoop      0 2015-06-11 19:00 /user
drwxrwxr-x  - hdfs hadoop      0 2015-06-11 18:49 /var
```

/tmp public writeable, used by many hadoop based application as temporary space.

/user all users home directory /user/\$USERNAME

/var public readable, used by many hadoop based application to store log files etc.

Working with HDFS

HDFS has file system shell

hadoop fs [commands]

The file system shell includes a set of command to work with hdfs.

Command are similar to common linux commands e.g.

>**hadoop fs -ls** #to list content of the default user directory.

>**hadoop fs -mkdir abc** #to make a directory in hdfs.

Getting Data in and out HDFS

hadoop fs -put local_file [path_in_HDFS]

Put a file in your local system into the HDFS

Each file would be stored in one or more “blocks”

The default block size is 128MB.

The block size used can be override by users

Hadoop fs -get path_in_hdfs [path_in_local]

Get a file from the hadoop cluster to the local file system.

Other File Shell commands

-stat returns stat information of a path

-cat/tail output to stdout

-setrep set replication factor

For a complete lists just do

hadoop fs

Or

visit <http://hadoop.apache.org/docs/r2.5.2/hadoop-project-dist/hadoop-common/FileSystemShell.html>

YARN

YARN: Yet Another Resource Manager

Managing computing resources within Hadoop cluster

All jobs should be submitted to yarn to run.

E.g. using either **yarn Jar** or **hadoop Jar**

When use other hadoop-supported application, please also specify YARN as resource manager. Such as SPARK.

YARN commands

Show cluster status

Help managing jobs running inside the Hadoop cluster.

YARN Commands

yarn application

-list to list applications submitted to YARN

default will show active/queued application

```
Total number of applications (application-types: [] and st
                                Application-ID      Application-Name
application_1431812610805_4832                           GS0108
```

-kill to kill application specified by application-ID.

-appStates/appTypes filter options

YARN Commands

yarn node

-list list of status of data nodes

Let us know if there is less than expected live data nodes.

yarn logs

dump logs of a finished application

-applicationID specific log from which application

-containerID specify log from which container

Running Hadoop Application

All Hadoop application can be run as a console command.

The basic format is like following:

*hadoop jar java_jar_name java_class_name
[parameters]*

The user can use –D to specify more Hadoop options.
e.g.

- D mapred.map.tasks #number of map instances to be generated.
- D mapred.reduce.tasks # number of reduce instances to be used.

Running Wordcount example code

The Hadoop distribution comes with an example jar which includes a set of exemplar map reduce code:

To run wordcount example from that jar file, you can run command like following:

```
hadoop jar /usr/lib/hadoop-mapreduce/hadoop-mapreduce-examples.jar \
wordcount \ #java class name to run
-D mapred.map.tasks=500 \ #number of mapper instance
-D mapred.reduce.tasks=256 \ # number of reducer instance
/tmp/data/enwiki-20120104-pages-articles.xml \ #input file on hdfs
wiki_wc \ #folder to store the output
```

Running Wordcount Example code

The on-screen output will show job running status

```
15/05/29 02:50:59 INFO client.RMProxy: Connecting to ResourceManager at name.rustler.tacc.utexas.edu/129.114.57.132:8032
15/05/29 02:51:01 INFO input.FileInputFormat: Total input paths to process : 1
15/05/29 02:51:02 INFO mapreduce.JobSubmitter: number of splits:266
15/05/29 02:51:02 INFO Configuration.deprecation: mapred.reduce.tasks is deprecated. Instead, use mapreduce.job.reduces
15/05/29 02:51:02 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
15/05/29 02:51:02 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1431812610805_3390
15/05/29 02:51:03 INFO impl.YarnClientImpl: Submitted application application_1431812610805_3390
15/05/29 02:51:03 INFO mapreduce.Job: The url to track the job: http://name.rustler.tacc.utexas.edu:8088/proxy/application_1431812610805_3390/
15/05/29 02:51:03 INFO mapreduce.Job: Running job: job_1431812610805_3390
15/05/29 02:51:03 INFO mapreduce.Job: Job job_1431812610805_3390 running in uber mode : false
15/05/29 02:51:09 INFO mapreduce.Job: map 0% reduce 0%
15/05/29 02:51:20 INFO mapreduce.Job: map 14% reduce 0%
15/05/29 02:51:21 INFO mapreduce.Job: map 24% reduce 0%
15/05/29 02:51:22 INFO mapreduce.Job: map 32% reduce 0%
15/05/29 02:51:23 INFO mapreduce.Job: map 45% reduce 0%
15/05/29 02:51:24 INFO mapreduce.Job: map 52% reduce 0%
15/05/29 02:51:25 INFO mapreduce.Job: map 57% reduce 0%
15/05/29 02:51:26 INFO mapreduce.Job: map 62% reduce 0%
15/05/29 02:51:27 INFO mapreduce.Job: map 65% reduce 0%
15/05/29 02:51:28 INFO mapreduce.Job: map 66% reduce 0%
15/05/29 02:51:30 INFO mapreduce.Job: map 67% reduce 0%
15/05/29 02:51:30 INFO mapreduce.Job: map 68% reduce 0%
15/05/29 02:51:30 INFO mapreduce.Job: map 72% reduce 0%
15/05/29 02:51:30 INFO mapreduce.Job: map 76% reduce 0%
15/05/29 02:51:41 INFO mapreduce.Job: map 81% reduce 0%
15/05/29 02:51:42 INFO mapreduce.Job: map 87% reduce 0%
15/05/29 02:51:43 INFO mapreduce.Job: map 93% reduce 0%
15/05/29 02:51:44 INFO mapreduce.Job: map 97% reduce 0%
15/05/29 02:51:45 INFO mapreduce.Job: map 98% reduce 0%
15/05/29 02:51:46 INFO mapreduce.Job: map 99% reduce 0%
15/05/29 02:51:47 INFO mapreduce.Job: map 100% reduce 0%
15/05/29 02:51:49 INFO mapreduce.Job: map 100% reduce 10%
15/05/29 02:51:50 INFO mapreduce.Job: map 100% reduce 24%
15/05/29 02:51:51 INFO mapreduce.Job: map 100% reduce 42%
15/05/29 02:51:52 INFO mapreduce.Job: map 100% reduce 76%
15/05/29 02:51:53 INFO mapreduce.Job: map 100% reduce 98%
15/05/29 02:51:54 INFO mapreduce.Job: map 100% reduce 100%
15/05/29 02:51:57 INFO mapreduce.Job: Job job_1431812610805_3390 completed successfully
15/05/29 02:51:58 INFO mapreduce.Job: Counters: 50
    File System Counters
        FILE: Number of bytes read=12501618427
        FILE: Number of bytes written=25053078456
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=35618393055
        HDFS: Number of bytes written=6489193532
        HDFS: Number of read operations=1566
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=512
    Job Counters
        Launched map tasks=266
        Launched reduce tasks=256
        Data-local map tasks=213
        Rack-local map tasks=53
        Total time spent by all maps in occupied slots (ms)=15587438
        Total time spent by all reduces in occupied slots (ms)=5572760
        Total time spent by all map tasks (ms)=7793719
        Total time spent by all reduce tasks (ms)=2786380
```

Parallelism with HDFS

The number of actual mapper created may be limited by the number of actual data blocks in the hdfs.

In the example, the input file is about ~35GB, with default block size as 128MB, the file is stored in 266 blocks in hdfs. All mappers may not be run at the same time, if there is not enough resources.

Each reducer will generate an output file independent from each other.

So 256 reducer would result 256 files in the output folder.

Running WordCount Example

Each output file is a text file as well,

Each line contains a word and its count,

You will notice for each output file, the word is sorted in alphabetic order

You can copy the file out of hdfs using command like

```
Hadoop fs –get /tmp/wiki_wc wiki_wc
```

Or you can view the content of each output using command like

```
Hadoop fs –cat /tmp/wiki_wc/part-r-00238 | more
```

Other Examples

There are more examples in the example.jar. You can use following command to get the list

```
hadoop jar /usr/lib/hadoop-mapreduce/hadoop-mapreduce-examples.jar
```

Here are a few examples:

grep: A map/reduce program that counts the matches of a regex in the input.

pi: A map/reduce program that estimates Pi using a quasi-Monte Carlo method.

terasort: sorting large set of random generated 100bytes data

sudoku: A sudoku solver.

Recap

- Hadoop and YARN
- Word Count Example – key value pair
 - Mapper (word, 1)
 - Reducer (word, count)
- Create a Hadoop Reservation
 - <https://portal.wrangler.tacc.utexas.edu/>
- Access Hadoop Reservation
 - VNC job – access Hadoop web UI etc.(TACC vis portal)
 - Idev job – interactive development/debugging
 - Batch job – submit bug-free code to Hadoop cluster

Try It Out

- Example code/data location

/work/00791/xwj/DMS/hadoop-training/

Copy to your home directory e.g.

> cp -r /work/00791/xwj/DMS/hadoop-training ~/hadoop

➤ cd ~/hadoop

Take a look of **exercise.txt** (<https://goo.gl/pUXO7g>)

Try exercise 1 and exercise 2 first

hadoop fs command syntax

hadoop fs -{ls | rm | cat| mkdir | put | get}

Using Hadoop with other programming language

Enabling MR jobs in other scripting language:
Python, Perl, R, C, etc...

User need to provide scripts/programs for Map and Reduce processing

The input/output format need to be compatible with key-value pair

Intermediate data are passed through stdin, stdout

A trade-off between convenience and performance

Hadoop Streaming API

```
hadoop jar /usr/lib/hadoop-mapreduce/hadoop-streaming.jar  
  -input /path/to/input/in/hdfs          #input file location  
  -output /path/to/output/in/hdfs       #output file location  
  -mapper map                          # mapper  
implementation  
  -reducer reduce                     #reduce  
implementation  
  -file map                           #location of the map code on local file  
system  
  -file reduce                        # location of the reduce code on local file  
system.
```

The map and reduce could be implemented in any programming language, even with bash script.

WordCount using Hadoop with bash

Putting together:

```
hadoop jar /usr/lib/hadoop-mapreduce/hadoop-streaming.jar \
-D mapred.map.tasks=512 \
-D mapred.reduce.tasks=256 \
-D stream.num.map.output.key.fields=1 \ # specify the position of the
key
-input /tmp/data/20news-all/alt.atheism \
-output wiki_wc_bash \
-mapper ./mapwc.sh -reducer ./reducewc.sh \
-file ./mapwc.sh -file ./reducewc.sh
```

WordCount using Hadoop with Python

```
hadoop jar $HADOOP_STREAMING \
-D mapred.map.tasks=4 -D mapred.reduce.tasks=2 \
-file mapper.py -mapper mapper.py \
-file reducer.py -reducer reducer.py \
-input /user/$USER/data/book.txt -output /user/$USER/
output-streaming-py
```

Hadoop vs Spark

Parameters	Spark	Hadoop
Data Storage	Spark stores data in-memory.	Hadoop stores data on disk.
Fault tolerance	Spark's data storage model, resilient distributed datasets (RDD) guarantees fault tolerance.	It uses replication to achieve fault tolerance.
Line of code	Apache Spark is project of 20,000 Line of code.	Hadoop 2.0 has 1,20,000 Line of code
Speed	It is Faster due to In-memory computation.	It is relatively slower than Spark.
OS Support	<ul style="list-style-type: none">• Linux• Windows• Mac OS	<ul style="list-style-type: none">• Linux
High level language	<ul style="list-style-type: none">• Scala• Python• Java• R	<ul style="list-style-type: none">• Java
Streaming data	Spark can be used to process as well as modify real-time data with Spark streaming.	With Hadoop Map-Reduce one can process batch of stored data.
Machine Learning	Spark has its own set of Machine learning libraries (<u>MLib</u>).	Hadoop requires interface with other Machine learning library. <u>Eg:</u> Apache Mahout.

Introduction of Zeppelin

- A web-based notebook that enables interactive data analytics and visualization.
- Multiple Language Backend



Start Zeppelin on Wrangler

```
login1.wrangler(5)$cd $DATA
```

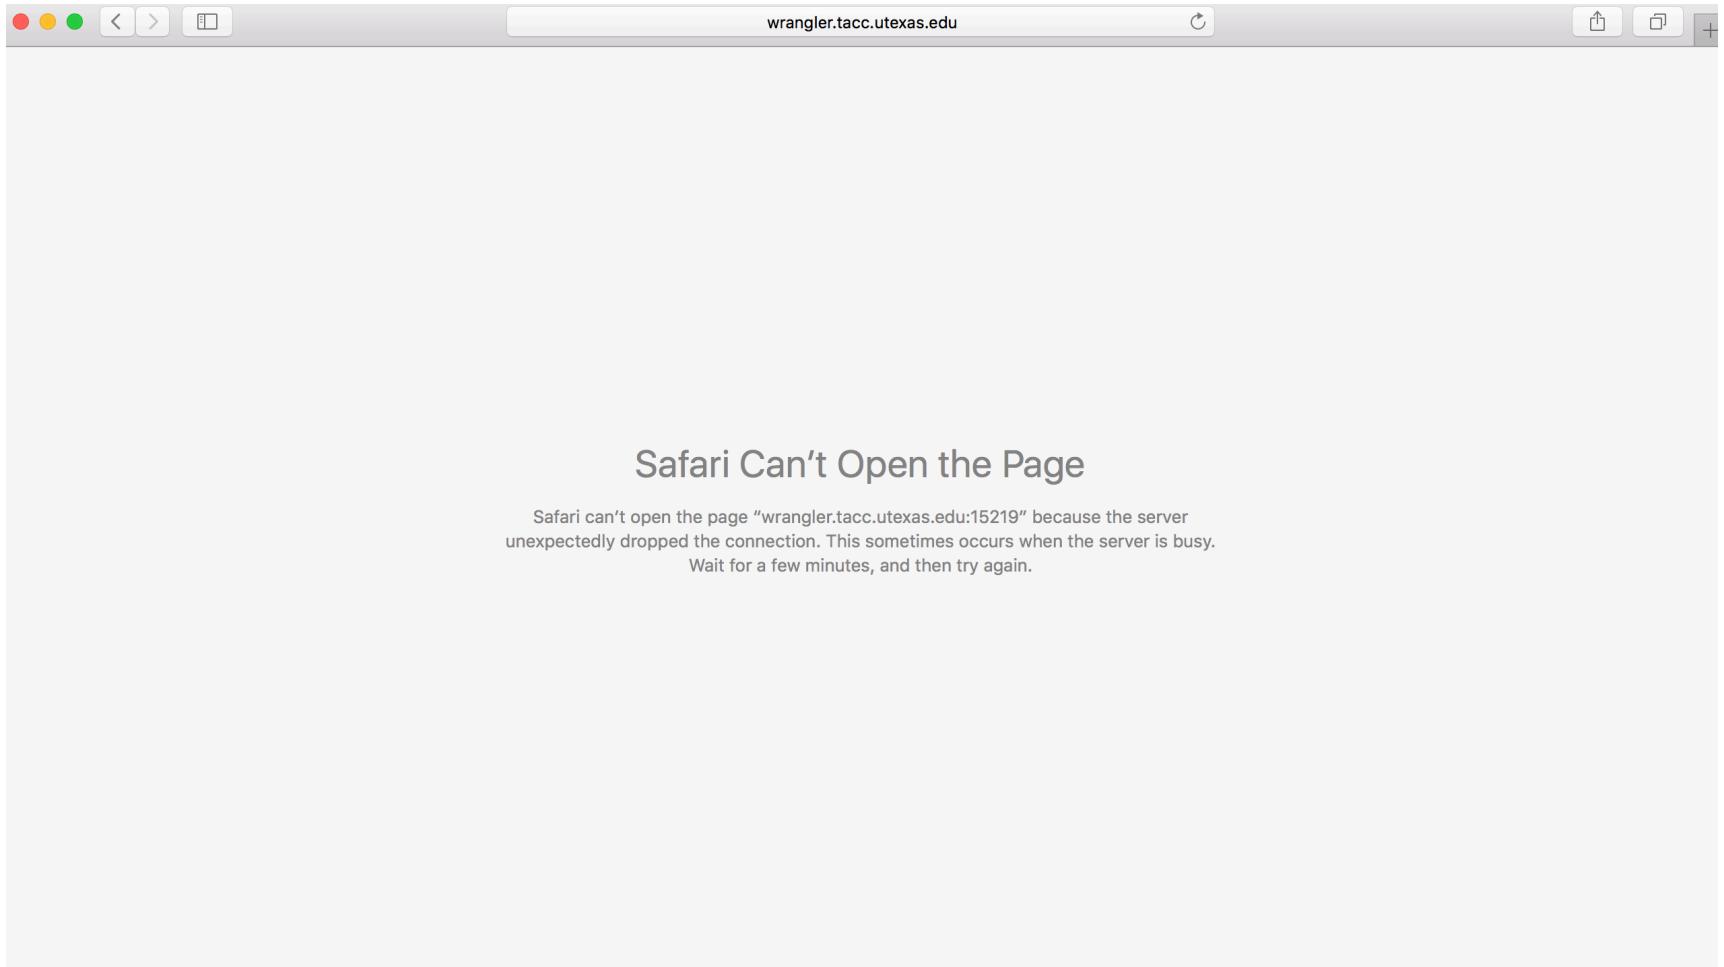
```
login1.wrangler(6)$cp /data/apps/.zeppelin/job.zeppelin .
```

```
login1.wrangler(7)$sbatch –  
reservation=hadoop+TRAINING-HPC+2186 job.zeppelin
```

```
login1.wrangler(15)$ tail zeppelin.out |tail -n 3  
Your applicatin is now running!  
Application UI is at http://wrangler.tacc.utexas.edu:15211  
Zeppelin username and password: user9062
```

Wait 10 minutes for Zeppelin UI to start, copy and paste
<http://wrangler.tacc.utexas.edu:XXXXX> to your web browser.

Use the username and password to login: userXXXX





Welcome to Zeppelin!

Zeppelin is web-based notebook that enables interactive data analytics.
You can make beautiful data-driven, interactive, collaborative document with SQL, code and even more!

Help

Get started with [Zeppelin documentation](#)

Community

Please feel free to help us to improve Zeppelin,
Any contribution are welcome!

- [Mailing list](#)
- [Issues tracking](#)
- [Github](#)

Wait 10 mins
until the red
light turns green

Login

Wrangler Zeppelin

wrangler.tacc.utexas.edu

Login

Welcome to Zeppelin!

Zeppelin is web-based notebook that enables interactive data analytics.
You can make beautiful data-driven, interactive, collaborative document with SQL, code and even more!

Help

Get started with [Zeppelin documentation](#)

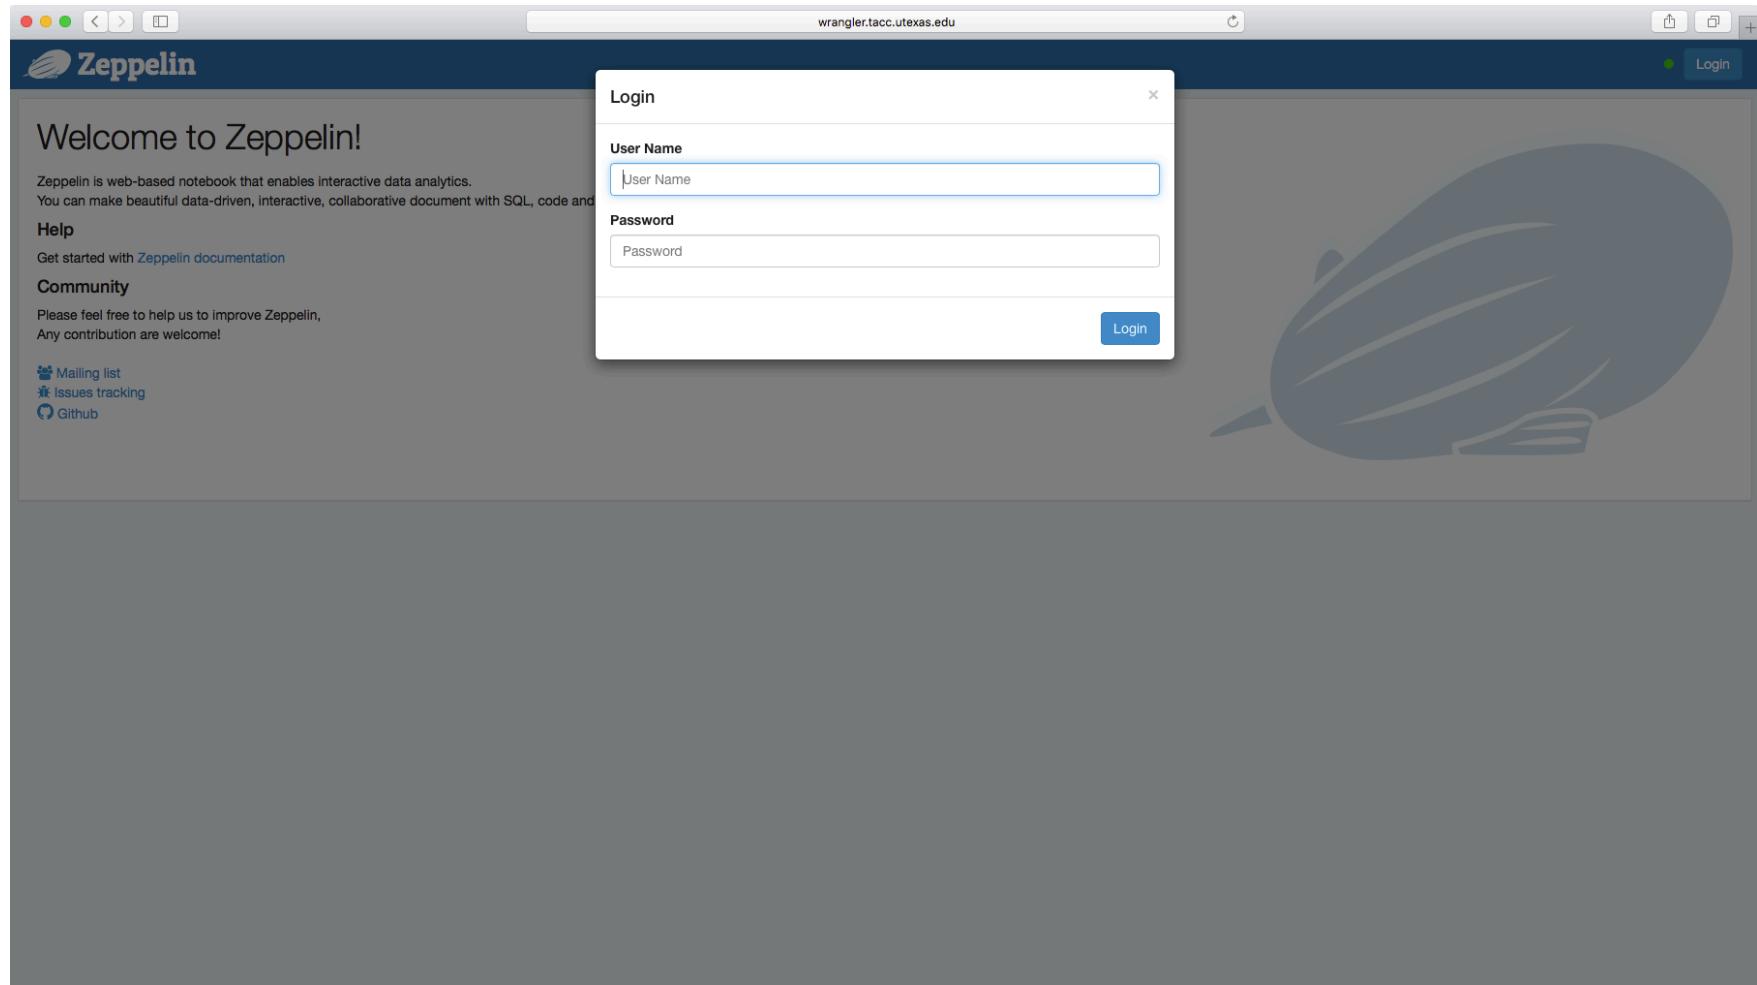
Community

Please feel free to help us to improve Zeppelin,
Any contribution are welcome!



-  Mailing list
-  Issues tracking
-  Github

Type in your username and password: userxxxx



wrangler.tacc.utexas.edu

Zeppelin Notebook Job Search your Notes user2181

Welcome to Zeppelin!

Zeppelin is web-based notebook that enables interactive data analytics.
You can make beautiful data-driven, interactive, collaborative document with SQL, code and even more!

Notebook

- [Import note](#)
- [Create new note](#)

Filter

- [test](#)
- [Zeppelin Tutorial](#)
- [Zeppelin_Demo](#)

Help

Get started with [Zeppelin documentation](#)

Community

Please feel free to help us to improve Zeppelin,
Any contribution are welcome!

- [Mailing list](#)
- [Issues tracking](#)
- [Github](#)



wrangler.tacc.utexas.edu

Zeppelin Notebook Job Search your Notes user2181

Welcome to Zeppelin!

Zeppelin is web-based notebook that enables interactive data analytics.
You can make beautiful data-driven, interactive, collaborative document with SQL, code and

Notebook

- Import note
- Create new note

Filter

- test
- Zeppelin Tutorial
- Zeppelin_Demo

Import new note

Import AS Note name

JSON file size cannot exceed 1 MB

Choose a JSON here Add from URL

The screenshot shows the Zeppelin web interface. A modal dialog box titled "Import new note" is open in the center. It has a field labeled "Import AS" with "Note name" and a note that "JSON file size cannot exceed 1 MB". There are two main options: "Choose a JSON here" represented by a cloud icon with an upward arrow, and "Add from URL" represented by a linked chain icon. The background of the page shows a blurred "Welcome to Zeppelin!" message and a sidebar with a "Notebook" section containing "Import note" and "Create new note" buttons, along with a "Filter" input field and a list of existing notebooks: "test", "Zeppelin Tutorial", and "Zeppelin_Demo". The top right corner shows a search bar "Search your Notes" and a user profile "user2181".

wrangler.tacc.utexas.edu

Welcome to Zeppelin

Zeppelin is web-based notebook that enables interactive data analysis. You can make beautiful data-driven, interactive, collaborative documents.

Notebook

Import note

Create new note

Create new note

Note Name

Default Interpreter

Use '/' to create folder

- ✓ spark
- md
- angular
- sh
- livy
- alluxio
- file
- psql
- flink
- python
- ignite
- lens
- cassandra
- kylinq
- elasticsearch
- jdbc
- hbase
- bq
- pig

Using Spark with Scala

```
%spark
val stopWords = sc.textFile("stopwords.txt")
val stopWordsSet = stopWords.collect.toSet
val stopWordSetBC = sc.broadcast(stopWordSet)

import org.apache.spark.sql.Row
//textFile.flatMap(_.toLowerCase.split(" ")).subtract(stopWords).take(100)
val wordCounts = textFile.flatMap(_.toLowerCase.split(" ")).filter(w => !stopWordSetBC.value.contains(w&&w!="")).map(word => (word, 1)).reduceByKey((a, b) => a + b)

val top50 = wordCounts.sortBy(_._2, ascending=false).map{case (word:String, count:Int) => {word + "\t" + count} }.take(50)
//top50.mkString("\n")
print("%table Word\t Count\n" + top50.mkString("\n"))

stopWords: org.apache.spark.rdd.RDD[String] = stopwords.txt MapPartitionsRDD[3] at textFile at <console>:28
stopWordSet: scala.collection.immutable.Set[String] = Set(serious, latterly, down, side, moreover, please, ourselves, behind, for, find, further, mill, due, any, wherein, across, twenty, name, this, in, move, itse", have, your, off, once, are, is, his, why, too, among, everyone, show, empty, already, nobody, less, am, hence, system, than, four, fire, anyhow, three, whereby, con, twelve, throughout, but, whether, below, co, mine, becomes, eleven, what, would, although, elsewhere, another, front, if, hereby, own, neither, bottom, up, etc, so, our, per, therein, must, beforehand, keep, do, all, him, had, somehow, re, onto, nor, every, herein, full, before, afterwards, somewhere, whither, else, namely, us, it, whereupon, two, thence, a, herse", sometimes, became, though, within, as, because...
stopWordSetBC: org.apache.spark.broadcast.Broadcast[scala.collection.immutable.Set[String]] = Broadcast(4)
import org.apache.spark.sql.Row
wordCounts: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[7] at reduceByKey at <console>:37
top50: Array[String] = Array(great 165, men 110, time 101, years 95, new 92, people 87, project 85, thousand 75, work 65, god 65, history 63, england 62, ancient 62, religious 61, long 58, city 57, human 40, gave 39, ne
arly 39, world. 39, a. 38, d. 38, days 37, received 37, hand 37, little 37, vast 37, place 36, british 36, [image: 36, modern 35)

settings ▾
```

Word	Count
time	110
people	92
work	65
england	63
long	58
gutenberg-tm	57
american	55
old	48
chapter	45
called	44
-	43
nations	42
world.	41
days	40
little	39
british	38
[image:]	36
modern	35

Took 1 min 57 sec. Last updated by anonymous at April 17 2017, 1:18:18 PM. (outdated)

Using Spark with Python

```
%pyspark
textFile = sc.textFile("book.txt")

stopWords = sc.textfile("stopwords.txt")

wordCounts = textFile.flatMap(lambda line: line.lower().split()).subtract(stopWords).map(lambda word: (word, 1)).reduceByKey(lambda a, b: a+b)
top50 = wordCounts.sortBy(lambda a: a[1], ascending=False).map(lambda x: x[0] + "\t" + str(x[1])).take(50)

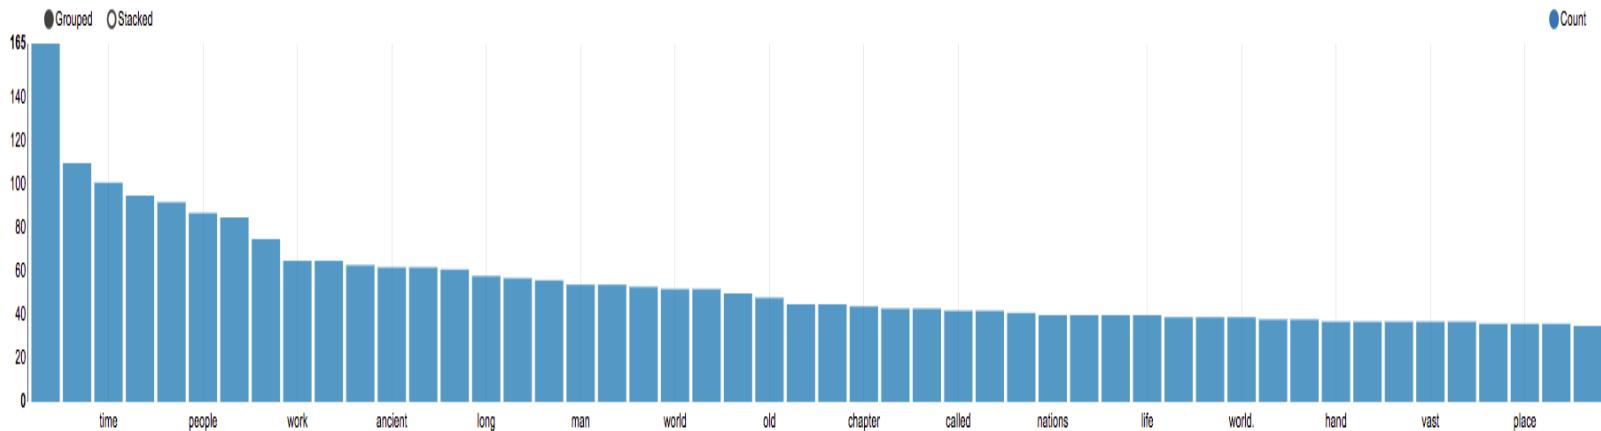
print("%table Word\t Count\n" + "\n".join(top50))
```

READY ▶ 🔍 ⌂ ⌂



settings▼

Count



Using Spark with R

SparkR interpreter setting(spark.executor.extraLibraryPath):

spark.executor.extraLibraryPath	/opt/apps/intel15/mvapich2_2_1/RstatsPackages/3.2.1/jags-3.4.0/lib64/JAGS/module-s-3:/opt/apps/intel15/mvapich2_2_1/RstatsPackages/3.2.1/jags-3.4.0/lib64:/opt/apps/intel15/mvapich2_2_1/RstatsPackages/3.2.1/proj-4.7.0/lib:/opt/apps/intel15/mvapich2_2_1/Rstats Packages/3.2.1/protobuf-4.7.1/lib:/opt/apps/intel15/mvapich2_2_1/RstatsPackages/3.2.1/gdal-1.9.2/lib:/opt/apps/intel15/mvapich2_2_1/RstatsPackages/3.2.1/nlopt-2.4.2/lib:/opt/apps/gcc/4.9.1/lib64:/opt/apps/gcc/4.9.1/lib:/opt/apps/intel15/mvapich2_2_1/Rstats/3.2.1/lib64/R/lib:/opt/apps/intel15/mvapich2/2.1/lib:/opt/apps/intel15/mvapich2_2_1/lib/shared:/opt/apps/intel15/composer_xe_2015.3.187/imprt/lib/intel64:/opt/apps/intel15/composer_xe_2015.3.187/ipp/lib/intel64:/opt/apps/intel15/composer_xe_2015.3.187/tbb/lib/intel64:/opt/apps/intel15/composer_xe_2015.3.187/tbb/lib/intel64/gcc4.4:/opt/apps/intel15/composer_xe_2015.3.187/compiler/lib/intel64
spark.executor.memory	
spark.r.command	/opt/apps/intel15/mvapich2_2_1/Rstats/3.2.1/bin/Rscript

```
%spark.r
##### Run a given function on a large dataset using dapply or dapplyCollect
# Convert waiting time from hours to seconds.
# Note that we can apply UDF to DataFrame.
schema <- structType(structField("eruptions", "double"), structField("waiting", "double"),
                      structField("waiting_secs", "double"))
df1 <- dapply(df, function(x) { x <- cbind(x, x$waiting * 60) }, schema)
head(collect(df1))

eruptions waiting waiting_secs
1      3.600      79      4740
2      1.800      54      3240
3      3.333      74      4440
4      2.283      62      3720
5      4.533      85      5100
6      2.883      55      3300
```

FINISHED ▶ ✎ 📄 ⚙

wrangler.tacc.utexas.edu

Zepplin Notebook Job Search your Notes user2181

Zepplin_Demo

%sh
echo \$LD_LIBRARY_PATH
copy and paste the below and set spark.executor.extraLibraryPath in the interpreter spark setting
/opt/apps/intel15/mvapich2_2_1/RstatsPackages/3.2.1/jags-3.4.0/lib64/JAGS/modules-3:/opt/apps/intel15/mvapich2_2_1/RstatsPackages/3.2.1/jags-3.4.0/lib64:/opt/apps/intel15/mvapich2_2_1/RstatsPackages/3.2.1/gdal-1.9.2/lib:/opt/apps/intel15/mvapich2_2_1/RstatsPackages/3.2.1/gcc-4.9.1/lib64:/opt/apps/gcc-4.9.1/lib:/opt/apps/intel15/mvapich2_2_1/Rstats/3.2.1/lib64/R/lib:/opt/apps/intel15/mvapich2_2.1/lib:/opt/apps/intel15/mvapich2_2.1/lib/shared:/opt/apps/intel/3.187/mpir/lib/intel64:/opt/apps/intel/15/composer_xe_2015.3.187/ipp/lib/intel64:/opt/apps/intel/15/composer_xe_2015.3.187/mkl/lib/intel64:/opt/apps/intel/15/composer_xe_2015.3.187/tbb/lib/intel5/composer_xe_2015.3.187/tbb/lib/intel64/gcc4.4:/opt/apps/intel/15/composer_xe_2015.3.187/compiler/lib/intel64

%spark.r
#detach("package:dplyr", unload=TRUE)

people <- read.df(sprintf("file:%s/examples/src/main/resources/people.json", Sys.getenv('SPARK_HOME')), "json")
head(people)
printSchema(people)

From Hive tables
sql("CREATE TABLE IF NOT EXISTS src (key INT, value STRING)")
input = sprintf("file:%s/examples/src/main/resources/kv1.txt", Sys.getenv('SPARK_HOME'))
sql(sprintf("LOAD DATA LOCAL INPATH '%s' INTO TABLE src", input))
Queries can be expressed in HiveQL.
results <- sql("FROM src SELECT key, value")
printSchema(results)
results is now a SparkDataFrame
dim(results)
results

schema <- structType(structField("key", "integer"), structField("value", "string"),
 structField("key2", "double"))

df1 <- dapply(results, function(x) { x <- cbind(x, x\$key * 2) }, schema)
head(df1)

age name
1 NA Michael
2 30 Andy
3 19 Justin
root
 |- age: long (nullable = true)
 |- name: string (nullable = true)
SparkDataFrame[]
SparkDataFrame[]
root
 |- key: integer (nullable = true)

About Zeppelin
Interpreter (highlighted)
Notebook Repos
Credential
Helium
Configuration
Logout

wrangler.tacc.utexas.edu

Interpreters

Manage interpreters settings. You can create / edit / remove settings. Note can bind / unbind these interpreter settings.

spark

Option

The interpreter will be instantiated **Globally** in **shared** process.

Connect to existing process
 Set permission

Properties

name	value
args	
master	yarn-client
spark.app.name	Zeppelin
spark.cores.max	
spark.executor.extraLibraryPath	/opt/apps/intel15/mvapich2_2_1/RstatsPackages/3.2.1/jags-3.4.0/lib64/JAGS/modules-3:/opt/apps/intel15/mvapich2_2_1/RstatsPackages/3.2.1/jags-3.4.0/lib64:/opt/apps/intel15/mvapich2_2_1/RstatsPackages/3.2.1/proj-4.7.0/lib:/opt/apps/intel15/mvapich2_2_1/RstatsPackages/3.2.1/protobuf-4.7.1/lib:/opt/apps/intel15/mvapich2_2_1/RstatsPackages/3.2.1/gdal-1.9.2/lib:/opt/apps/intel15/mvapich2_2_1/RstatsPackages/3.2.1/nlopt-2.4.2/lib:/opt/apps/gcc-4.9.1/lib64:/opt/apps/gcc-4.9.1/lib:/opt/apps/intel15/mvapich2_2_1/Rstats/3.2.1/lib64/R/lib:/opt/apps/intel15/mvapich2/2.1/lib:/opt/apps/intel15/mvapich2/2.1/lib/shared:/opt/apps/intel/15/composer_xe_2015.3.187/mpir/lib/intel64:/opt/apps/intel/15/composer_xe_2015.3.187/ipp/lib/intel64:/opt/apps/intel/15/composer_xe_2015.3.187/mkl/lib/intel64:/opt/apps/intel/15/composer_xe_2015.3.187/tbb/lib/intel64:/opt/apps/intel/15/composer_xe_2015.3.187/tbb/lib/intel64/gcc4.4:/opt/apps/intel/15/composer_xe_2015.3.187/compiler/lib/intel64
spark.executor.memory	
spark.r.command	/opt/apps/intel15/mvapich2_2_1/Rstats/3.2.1/bin/Rscript
zeppelin.R.cmd	R
zeppelin.R.image.width	100%

[spark ui](#) [edit](#) [restart](#) [remove](#)

Zeppelin Notebook Job wrangler.tacc.utexas.edu Search your Notes user2181

spark.r.command	/opt/apps/intel15/mvapich2_2_1/Rstats/3.2.1/bin/Rscript
zeppelin.R.cmd	R
zeppelin.R.image.width	100%
zeppelin.R.knitr	true
zeppelin.R.render.options	out.format = 'html', comment = NA, echo = FALSE, results = 'asis', message = F, warning = F
zeppelin.dep.additionalRemoteRepository	spark-packages,http://dl.bintray.com/spark-packages/maven,false;
zeppelin.dep.localrepo	local-repo
zeppelin.interpreter.localRepo	/data/03076/rhuang/.zeppelin/local-repo/2CDA15AM4
zeppelin.pyspark.python	python
zeppelin.spark.concurrentSQL	false
zeppelin.spark.importImplicit	true
zeppelin.spark.maxResult	1000
zeppelin.spark.printREPOOutput	true
zeppelin.spark.sql.stacktrace	false
zeppelin.spark.useHiveContext	true

Dependencies
These dependencies will be added to classpath when interpreter process starts.

artifact	exclude	action
groupId:artifactId:version or local file path	(Optional) comma separated groupId:artifactId list	+

Save Cancel

Demo - Zeppelin

Note:

- Clear browser data (cached images and files) to avoid loading the cached image of Zeppelin UI.
- To change default setting in the Zeppelin's interpreter:
 - spark.executor.instances: 12
 - spark.executor.memory: 2g
 - spark.driver.memory: 4g

Recap

- **Hadoop vs. Spark**
 - Spark: in-memory, faster, interactive development, multi-languages (Scala, Java, Python, R)
- **Using Hadoop with other languages**
 - Also called Hadoop streaming – confusing naming
 - Hard to debug
- **Zeppelin**
 - Interactive multi-interpreters notebook with great visualization support
- **Using Spark with Scala, Python and R**
 - SparkR api quite different from Scala and Python

Try It Out

Exercise 3 shows an example of compile from user's java code to Hadoop application and run it.

Exercise 4 is about hadoop streaming

Can you write a word count example with your favorite programming language?

Exemplar code for Python, R and bash are available at the corresponding directory.

Questions to think?

Where/how did we specify parallelisms?

Exercise 5 is an example of a completed workflow for using Mahout for document clustering.

More Questions

Data Group @ TACC
data@tacc.utexas.edu

Weijia Xu
xwj@tacc.utexas.edu

Ruihu Huang
rhuang@tacc.utexas.edu

Thank You!

Misc.

The reservation name is hadoop+TRAINING-HPC+2186

Exemplar data, code, and practice command suggestions
are also in **/work/00791/xwj/DMS/hadoop-training**

idev session command

```
idev -r hadoop+TRAINING-HPC+2186 -m 240 -p hadoop
```

Vis portal vis.tacc.utexas.edu

Refer to exercise text or <https://goo.gl/pUXO7g>